


Fundamentals of Data Engineering

The Ultimate Trip.



Today

2

Welcome |
| Course Information |
| Data Modeling |
| Memory |
| File |
| Database

Welcome | Course Information | Data Modeling | Memory | File | Database

Course Information

3

Title	Fundamentals of Data Engineering
<u>Code</u>	COE848
Term	Winter 2019
Time	Tuesday 09AM-12PM
Location	KHE225

Lecturer

4

Name

Hossein Fani

Education

Ph.D. Candidate

Faculty of Computer Science

University of New Brunswick, NB., Canada

Office

ENG324

Hour

Tuesday 12PM-01PM

Email

hossein.fani@ryerson.ca

Grading

5

Labs and Final Project

35%

Midterm Exam (W08: Feb. 26, 2019)

25%

Final Exam

40%

Laboratory

6

Location

ENG411

Assignments

Steps of one final project

Deadlines

7 days

Late submission receives none!


Laboratory × Deadline

7

Lab#	No. of Weeks	Due Date
Lab 1	1	Jan. 23
Lab 2	1	Jan. 30
Lab 3	2	Feb. 13
Lab 4	2	Mar. 06
Lab 5	2	Mar. 20
Lab 6	2	Apr. 03
All sections have same submission deadline, i.e., Thursdays Midnight		

Laboratory × TAs

8



Chengliang Huang
c28huang@ryerson.ca



Raman Deep Singh
raman.singh@ryerson.ca

Laboratory × Time

9

	Monday	Tuesday	Wednesday	Thursday
09:00				
10:00		C: Hossein KHE225	L3: Raman ENG411	
11:00	L4: Chengliang ENG411			
12:00		O: Hossein ENG324		L2: Raman ENG411
13:00				
14:00	L1: Chengliang ENG411			
15:00				
L: Lab C: Class O: Office Hour				

Book

10

Title

Database Systems:
The Complete Book

Authors

Hector Garcia-Molina
Jeffrey D. Ullman
Jennifer D. Widom

Publisher


Prentice Hall

Edition

2nd Edition, 2008

ISBN

0131873253



Homepage (D2L)

11

Course Outline

Done

Lab Guide

Done

Lab Deadlines

Done

Slides

End of Each Week

Announcements

Check D2L for course updates & information!

Homepage (D2L) Tour

12

Running Example

13

Title

Moviesion

Use Case

Similar Movie Recommendation

Test Case

Input: "2001: A Space Odyssey"

Output: 1. "Interstellar"

2. "Her"

3. "Metropolis"

4. ...

Programming Language

14

Name

Java (why?)

IDE

NetBeans

Code Convention

www.oracle.com/technetwork/java/codeconventions-150003.pdf

Naming Convention

www.oracle.com/technetwork/java/codeconventions-135099.html

Programming Language × Naming

15

Package

all.lower.case

```
package ca.ryerson.ee.coe848
```

Class

Singular Noun, PascalCase

```
public class Movie{}  
public class MovieProductionCompany{}
```

Variable

Noun, camelCase

```
Movie oscarBestPicture = new Movie();  
ArrayList<Movie> oscarWinners = new ArrayList<>();
```

Method

Verb, camelCase

```
public void addMovie(Movie newMovie);
```

Constant

ALL UPPER CASE

```
final int MAX_MOVIE_TITLE_WIDTH = 255;
```

Attendance

16

Policy

KINDLY DO NOT DISTRACT OTHERS!

Phone

Off

Silent

Laptop

Closed

Far End of Each Row

Break

10:30 - 10:45

Others

www.ryerson.ca/studentguide/learning_experience/

The background is a dark, star-filled space. In the upper center, a satellite or space station is visible. A bright, glowing light source, possibly the sun, is positioned in the middle ground, creating a lens flare effect. The overall scene is framed by a dark, triangular shape at the bottom, suggesting a view from a spacecraft or a deep space environment.

2019: A DATA ODYSSEY

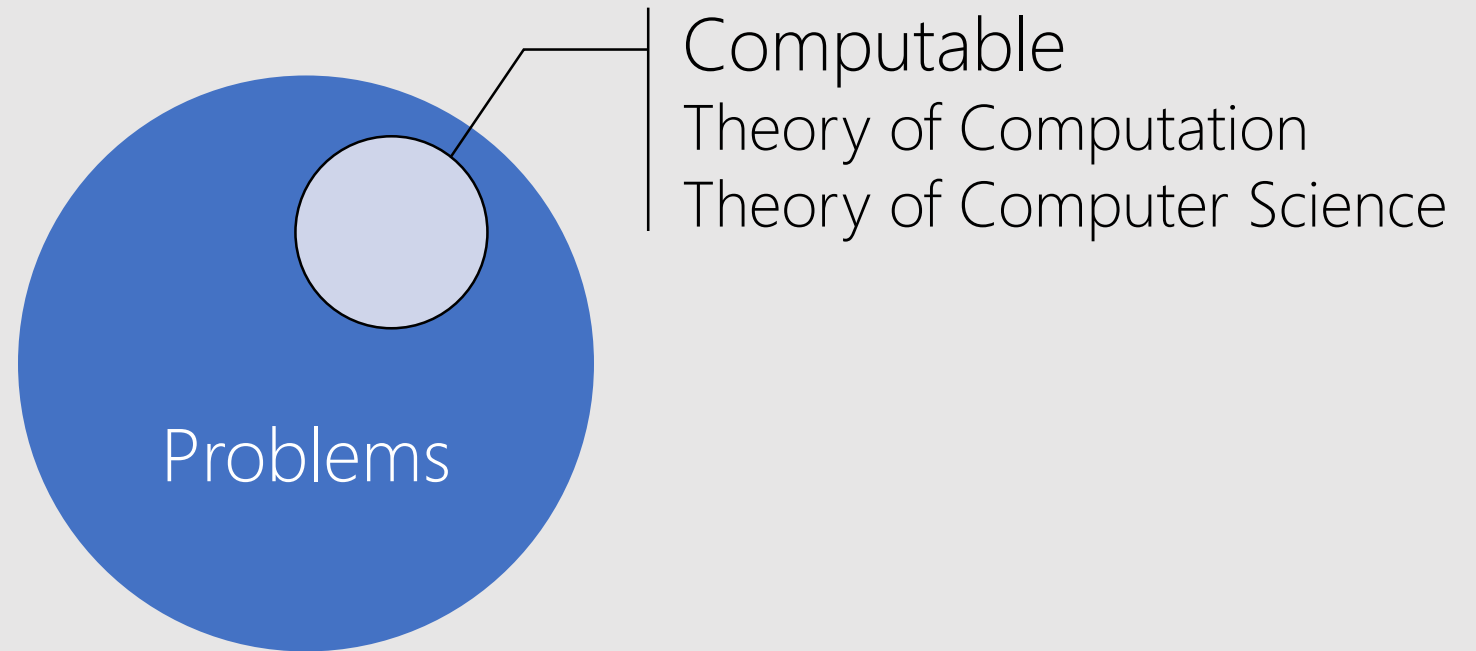
2019: A Data Odyssey

18




2019: A Data Odyssey

19



2019: A Data Odyssey


20



2019: A Data Odyssey

21

There is interaction between data engineers and algorithm designers.




Data gets bigger!
Algorithms get slower!

Better Data Modeling → Faster Algorithm
e.g., ?

2019: A Data Odyssey

21

There is interaction between data engineers and algorithm designers.




Data gets bigger!
Algorithms get slower!

Better Data Modeling → Faster Algorithm
e.g., Sorting in Arrays vs. Trees
e.g., Searching in Arrays vs. Trees
e.g., Sparse Matrix in Arrays vs. Linked Lists

2019: A Data Odyssey × Academy

23

Algorithm Design
Algorithm Analysis
Artificial Intelligence (AI)
Machine Learning
Data Mining




Data Structure (Memory)
File Structure (File Systems)
Database Management Systems
Data Warehouse
Big Data
Cloud

2019: A Data Odyssey × Academy

24

Algorithm Design
Algorithm Analysis
Artificial Intelligence (AI)
Machine Learning
Data Mining



Data Structure (Memory)
File Structure (File Systems)
Database Management Systems
Data Warehouse
Big Data
Cloud

2019: A Data Odyssey × Real World

25

Data Modeling: Real World Entity

Conceptual Level |

| Logical Level |

| Physical Level

Computable Entity

Data Modeling × Conceptual Level

26

1. Identify real world entities, attributes, relationships
2. Create Schema

Data Modeling × Conceptual Level


27

Entities for a movie recommender system:



Movie



Director



Company



Data Modeling × Conceptual Level

28

Movie Attributes:

Title


Genre

Language

RunningTime

Poster

...



Data Modeling × Conceptual Level

29

Movie Attributes' Type (More Optional Details):

Title (`string`)


Genre (`string`)

Language (`string`)

RunningTime (`decimal`)

Poster (?)

...



Data Modeling × Conceptual Level

30

Movie Attributes' Type (More Optional Details):

Title (`string`)


Genre (`string`)

Language (`string`)

RunningTime (`decimal`)

Poster (`byte[]`)

...



Data Modeling × Conceptual Level

31

Director Attributes:


Name (`string`)

DateOfBirth (`date`)

PlaceOfBirth (`string`)

Photo (`byte[]`)

...



Data Modeling × Conceptual Level

32

Company Attributes:

Name (`string`)

Address (`string`)

DateOfEstablishment (`date`)

Logo (`byte[]`)

...




Data Modeling × Conceptual Level

33

Movie Schema	Movie Data <u>Instance</u>
Title (<i>string</i>)	<i>2001: A Space Odyssey</i>
Genre (<i>string</i>)	<i>Sci-fi</i>
Language (<i>string</i>)	<i>English</i>
RunningTime (<i>decimal</i>)	<i>142</i>
Poster (<i>byte[]</i>)	



Data Modeling × Conceptual Level

34

Movie Schema	Movie Data <u>Instance</u>
Title (<i>string</i>)	<i>Rosemary's Baby</i>
Genre (<i>string</i>)	<i>Horror</i>
Language (<i>string</i>)	<i>English</i>
RunningTime (<i>decimal</i>)	<i>136</i>
Poster (<i>byte[]</i>)	


Data Modeling × Conceptual Level

35

Movie Schema	Title (<code>string</code>)	Genre (<code>string</code>)	Language (<code>string</code>)	RunningTime (<code>decimal</code>)	Poster (<code>byte[]</code>)
Movie Instance 1	<i>2001: A Space Odyssey</i>	<i>Sci-fi</i>	<i>English</i>	<i>142</i>	
Movie Instance 2	<i>Rosemary's Baby</i>	<i>Horror</i>	<i>English</i>	<i>136</i>	



Data Modeling × Conceptual Level

36

Director Schema	Name (<code>string</code>)	DateOfBirth (<code>date</code>)	PlaceOfBirth (<code>date</code>)	Photo (<code>byte[]</code>)
Director Instance 1	<i>Stanley Kubrick</i>	<i>1949</i>	<i>USA</i>	
Director Instance 2	<i>Roman Polanski</i>	<i>1933</i>	<i>Poland</i>	

Data Modeling × Conceptual Level

37

Company Schema	Name (string)	Address (string)	DateOfEstablishment (date)	Logo (byte[])
Company Instance 1	<i>Metro-Goldwyn-Mayer Studios Inc.</i>	<i>Beverly Hills, CA, USA</i>	1924	
Company Instance 2	<i>Warner Bros</i>	null	1923	


Data Modeling × Conceptual Level

38

Data Schema	Data Instance
Data about Data: Meta-data	<i>Actual Data</i>
Defined at Setup Time	<i>Inserted at Running Time</i>
Rarely Change	<i>Rapidly Change</i>
	<i>Must Conform to Schema</i>


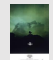
Data Modeling × Conceptual Level



39





Data Modeling × Conceptual Level

40

Movie Schema	Title (string)	Genre (string)	Language (string)	RunningTime (decimal)	Poster (byte[])
Movie Instance 1	<i>2001: A Space Odyssey</i>	<i>Sci-fi</i>	<i>English</i>	142	
Movie Instance 2	<i>Rosemary's Baby</i>	<i>Horror</i>	<i>English</i>	136	

Director Schema	Name (string)	DateOfBirth (date)	PlaceOfBirth (date)	Photo (byte[])
Director Instance 1	<i>Stanley Kubrick</i>	1949	USA	
Director Instance 2	<i>Roman Polanski</i>	1933	Poland	

Company Schema	Name (string)	Address (string)	DateOfEstablishment (date)	Logo (byte[])
Company Instance 1	<i>Metro-Goldwyn-Mayer Studios Inc.</i>	<i>Beverly Hills, CA, USA</i>	1924	
Company Instance 2	<i>Warner Bros</i>	null	1923	

Data Modeling × Conceptual Level

41

Movie Schema	Title (string)	Genre (string)	Language (string)	RunningTime (decimal)	Poster (byte[])
Movie Instance 1	2001: A Space Odyssey	Sci-fi	English	142	
Movie Instance 2	Rosemary's Baby	Horror	English	136	
Director Schema	Name (string)	DateOfBirth (date)	PlaceOfBirth (date)	Photo (byte[])	
Director Instance 1	Stanley Kubrick	1949	USA		
Director Instance 2	Roman Polanski	1933	Poland		
Company Schema	Name (string)	Address (string)	DateOfEstablishment (date)	Logo (byte[])	
Company Instance 1	Metro-Goldwyn-Mayer Studios Inc.	Beverly Hills, CA, USA	1924		
Company Instance 2	Warner Bros	null	1923		

No Relationship!

Movie × Director × Company

Data Modeling × Conceptual Level




42

Movie × Director × Company Relationships:

Director <u>makes</u> Movie		Movie is made by Director
Company <u>distributes</u> Movie		Movie is distributed by Company
Company <u>budgets</u> Movie		Movie is budgeted by Company







Data Modeling × Conceptual Level

43

Movie Schema	Title (string)	Genre (string)	Language (string)	RunningTime (decimal)	Poster (byte[])	DirectedBy (Director Schema)			
						Name (string)	DateOfBirth (date)	PlaceOfBirth (date)	Photo (byte[])
Movie Instance 1	<i>2001: A Space Odyssey</i>	<i>Sci-fi</i>	<i>English</i>	<i>142</i>		<i>Stanley Kubrick</i>	<i>1949</i>	<i>USA</i>	
Movie Instance 2	<i>Rosemary's Baby</i>	<i>Horror</i>	<i>English</i>	<i>136</i>		<i>Roman Polanski</i>	<i>1933</i>	<i>Poland</i>	







Data Modeling × Conceptual Level

44

Movie Schema	Title (string)	Genre (string)	Language (string)	RunningTime (decimal)	Poster (byte[])	DirectedBy (Director Schema)			
						Name (string)	DateOfBirth (date)	PlaceOfBirth (date)	Photo (byte[])
Movie Instance 1	2001: A Space Odyssey	Sci-fi	English	142		Stanley Kubrick	1949	USA	
Movie Instance 2	Rosemary's Baby	Horror	English	136		Roman Polanski	1933	Poland	
Movie Instance 3	A Clockwork Orange	Sci-fi	English	136		Stanley Kubrick	1949	USA	

Data Modeling × Conceptual Level

45

Movie Schema	Title (string)	Genre (string)	Language (string)	RunningTime (decimal)	Poster (byte[])	DirectedBy (Director Schema)			
						Name (string)	DateOfBirth (date)	PlaceOfBirth (date)	Photo (byte[])
Movie Instance 1	2001: A Space Odyssey	Sci-fi	English	142		Stanley Kubrick	1949	USA	
Movie Instance 2	Rosemary's Baby	Horror	English	136		Roman Polanski	1933	Poland	
Movie Instance 3	A Clockwork Orange	Sci-fi	English	136		Stanley Kubrick	1949	USA	

Redundancy!

Not All Data Models at Conceptual Level are Efficient.
Not All Conceptual Models are Efficient.
Not All Data Models are Efficient.

Data Modeling × Conceptual Level

46

Practice1: Data Modeling for Ryerson Library at Conceptual Level

Entities: ?

Attributes: ?

Relationships: ?

Data Modeling × Conceptual Level

47

Practice1: Data Modeling for Ryerson Library at Conceptual Level

Entities: Book, Staff, Person, ...

Attributes:

Book: ISBN, Title, ...

Staff: Name, Salary, Rank, ...

Student: StudentId, Name, ...

Relationships: Student borrows Book from Staff

Data Modeling × Conceptual Level

48

Practice2: Data Modeling for Ryerson University at Conceptual Level

Entities: ?

Attributes: ?

Relationships: ?

2019: A Data Odyssey × Real World

49

Data Modeling: Real World Entity

Conceptual Level |

| Logical Level |

| Physical Level

Computable Entity

Data Modeling × Logical Level

50

1. How entities, attributes, relationships should be represented.
2. Update Schema

Data Modeling × Logical Level



51

We already have seen one representation, haven't we?

Data Modeling × Logical Level

52

We already have seen one representation, haven't we? Tabular!

Movie Schema	Title (<code>string</code>)	Genre (<code>string</code>)	Language (<code>string</code>)	RunningTime (<code>decimal</code>)	Poster (<code>byte[]</code>)
Movie Instance 1	<i>2001: A Space Odyssey</i>	<i>Sci-fi</i>	<i>English</i>	142	
Movie Instance 2	<i>Rosemary's Baby</i>	<i>Horror</i>	<i>English</i>	136	

Data Modeling × Logical Level

53

There are other representations as well.

There are other logical models as well.

There are other data models at logical level as well.

- ~1960: Object Oriented

- 1969: Relational: Mathematical Relationships

- 1996: XML

Data Modeling × Logical Level

54

There are other representations as well.

There are other logical models as well.

There are other data models at logical level as well.

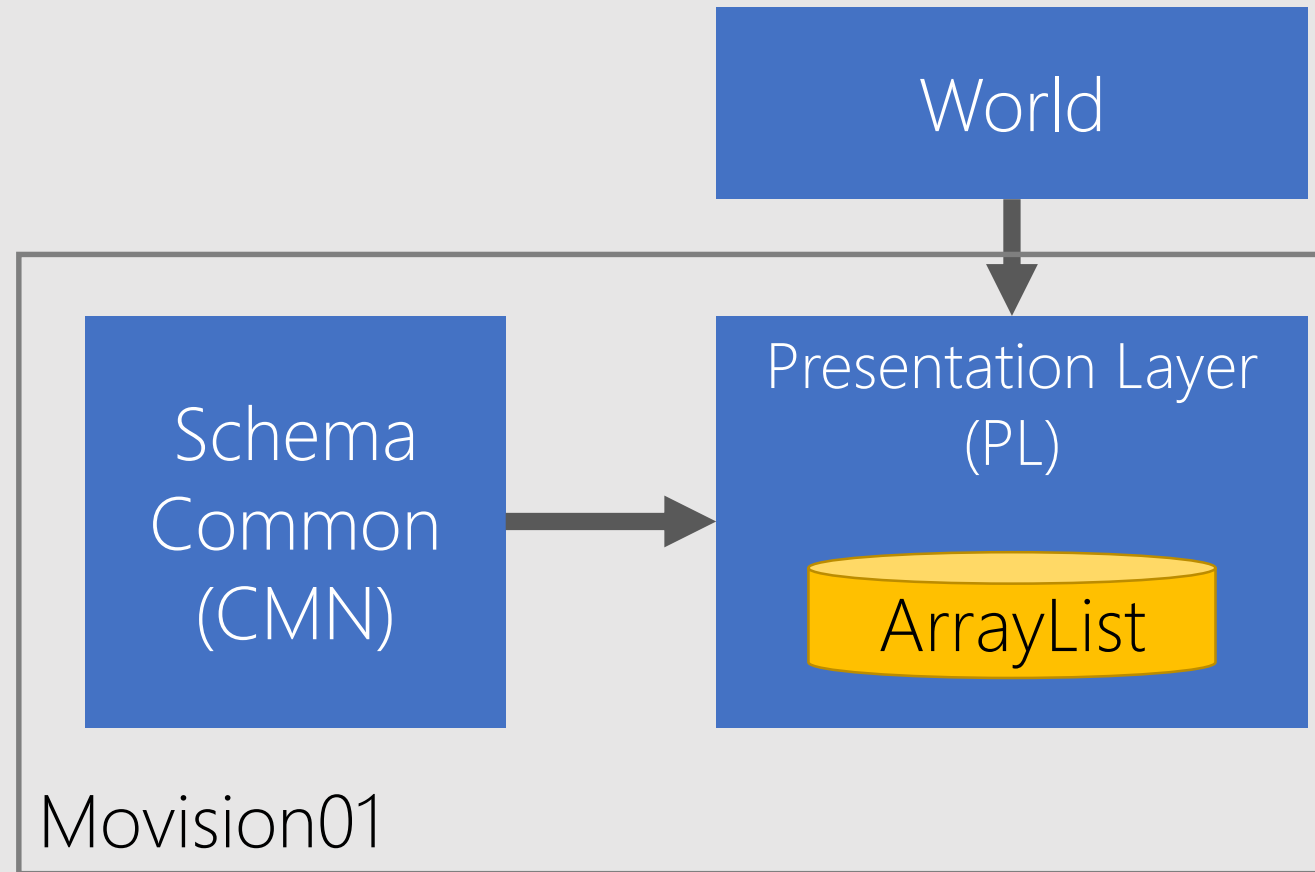
~1960: Object Oriented

1969: Relational: Mathematical Relation

1996: XML

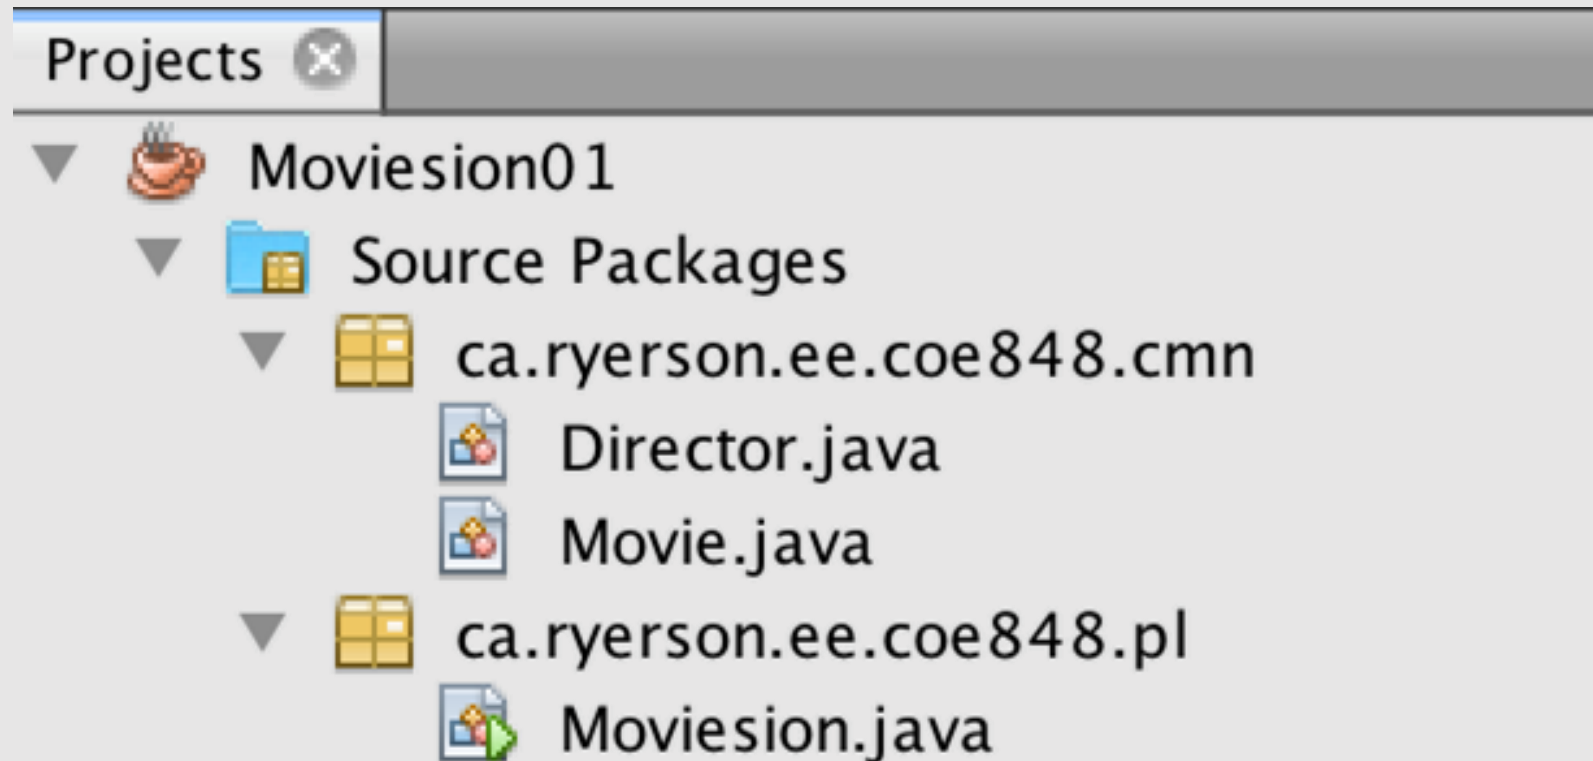
Logical Level × Object Oriented

55



Logical Level × Object Oriented

56



Logical Level × Object Oriented

57

Conceptual Model	Transformation	Logical Model by OO
Data Schema	DDL	<i>Class</i>
Data Instance	DML	<i>Object</i>

Data Definition Language (DDL): Commands for setting up schema of data

Data Manipulation Language (DML): Commands to manipulate data, also called "query language"

Logical Level × Object Oriented

58

Conceptual Model	Logical Model by OO
Movie Schema	<code>public class Movie {</code>
Title	String <code>title;</code>
Genre	String <code>genre;</code>
Language	String <code>language;</code>
RunningTime	Decimal <code>runningTime;</code>
Poster	Byte[] <code>poster;</code>
...	<code>}</code>

Logical Level × Object Oriented


59

Conceptual Model

Logical Model by OO

Movie Instance

Movie Object

Title	Genre	Language	RunningTime	Poster
2001: A Space Odyssey	Sci-fi	English	142	

```
Movie newMovie = new Movie();  
newMovie.setName("2001: A Space Odyssey");  
newMovie.setGenre("Sci-fi");  
newMovie.setLanguage("English");  
newMovie.setRunningTime(142);  
newMovie.setPoster(?);
```

Logical Level × Object Oriented

60

Pros

Time | Random Access Memory (RAM) | Fast

Variety of Data Structures

Array, Linked List, Stack, Queue, Tree, HashTable, ...

Easy DML in Memory

SELECT: `E element = ArrayList<E>.get(int index);`

INSERT: `ArrayList<E>.add(int index, E element);`

DELETE: `ArrayList<E>.remove(int index);`

UPDATE: `ArrayList<E>.set(int index, E element);`

Any Computation

Sorting, Searching, ...

Cons

Space Far Too Small

Expensive

NOT DURABLE

Volatility | Transient | Short-term Retention

Portability

2019: A Data Odyssey × Real World

61

Data Modeling: Real World Entity

Conceptual Level |

| Logical Level |

| Physical Level

Computable Entity

Data Modeling × Physical Level

62

1. Where and how are the data instances stored in physical storage.

Data Modeling × Physical Level

63

We already have seen one physical storage, haven't we?

Data Modeling × Physical Level

64

There are other physical storages as well.

There are other physical models as well.

There are other data models at physical level as well.

File

Database

Cloud

Data Modeling × Physical Level

65

There are other physical storages as well.

There are other physical models as well.

There are other data models at physical level as well.

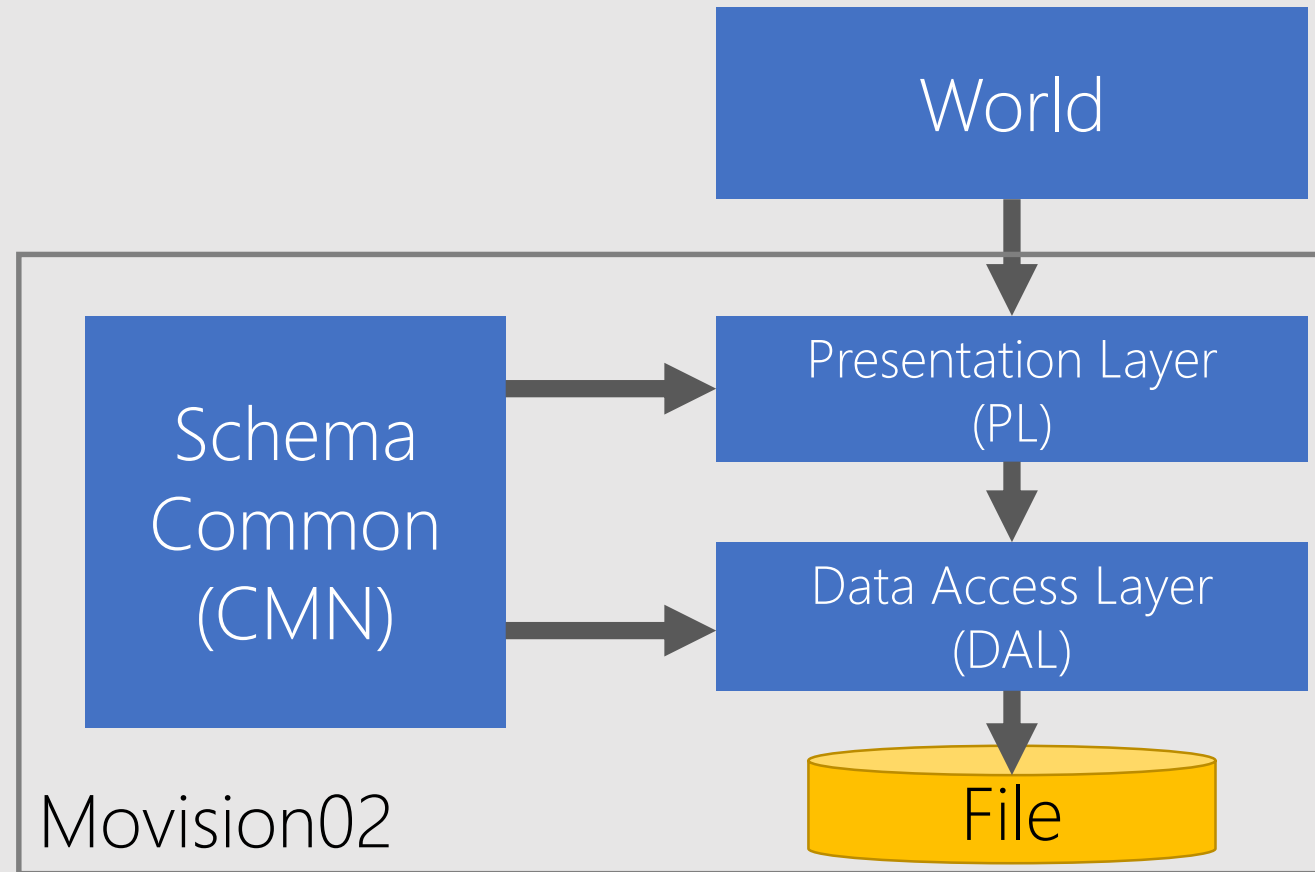
File

Database

Cloud

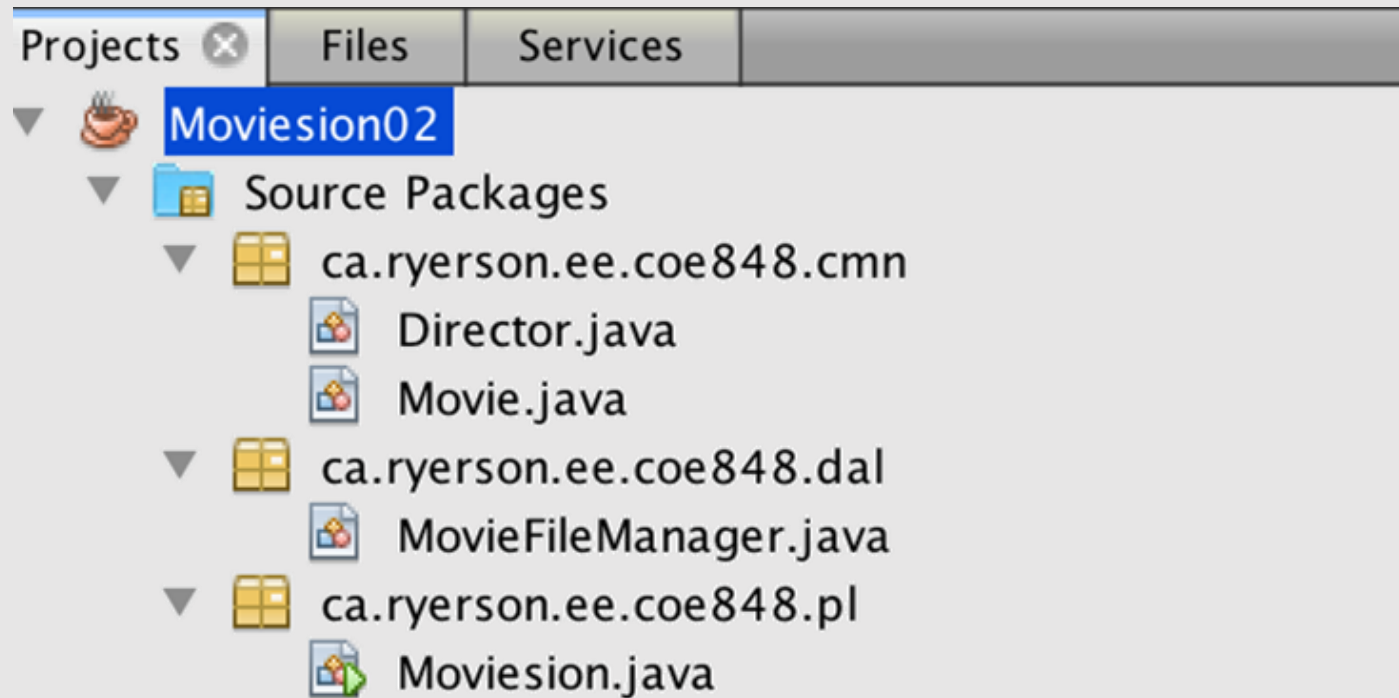
Physical Level × File

66



Physical Level × File

67



Physical Level × File

68

Conceptual Model	Logical Model by OO	Physical Model by File
Data Schema	Class	<i>CSV File: nothing</i> <i>JSON File: nothing</i> <u><i>XML File: XSD</i></u>
Data Instance	Object	<i>Record</i>

2019: A Data Odyssey × Real World

69

Real World Entity

Conceptual Level | Instance

| Logical Level | Object

| Physical Level | Record

Computable Entity

Physical Level × File

70

Pros

Cons

Space

Time

*Hard Disk Drive (HDD)
Sequential Access Method
Electromechanical*

DURABLE

Persistency | Long-term Retention

Not Easy DML

DELETE | UPDATE

Portable

No ACID Properties

Solid State Drive (SSD) | Random Access Memory | Expensive

Physical Level × File × Cons

71

Not Easy DML

DELETE

Strategy#1: Create a new file without the deleted record.

Strategy#2: Shift all following over the deleted record.

Strategy#3: Logical Deletion | Invalidate Record

UPDATE

Strategy#1: Replace bytes of old value with bytes of new value

Strategy#2: Delete Record + Insert Record

Physical Level × File × ACID

72

Transaction

A group one or more operations (DDL|DML) into a single unit of work.

```
BEGIN TRANSACTION
```

```
  INSERT Movie
```

```
  INSERT Movie's Director IF NOT EXIST
```

```
  INSERT Movie's distributor Company IF NOT EXIST
```

```
END TRANSACTION
```


Physical Level × File × ACID

73

Atomicity

All-or-nothing execution of transaction

```
BEGIN TRANSACTION
```

```
  INSERT Movie
```

```
  INSERT Movie's Director IF NOT EXIST
```

```
  INSERT Movie's distributor Company IF NOT EXIST
```

```
END TRANSACTION
```

Transaction Committed.

3 records affected.

Physical Level × File × ACID

74

Atomicity

All-or-nothing execution of transaction

BEGIN TRANSACTION

INSERT *Movie*

Error!

Partial Execution of Transaction!

ROLLBACK any changes, i.e., *Movie* record!

Physical Level × File × ACID

75

Atomicity

All-or-nothing execution of transaction

Another example in online banking?

Physical Level × File × ACID

76

Atomicity

All-or-nothing execution of transaction

BEGIN TRANSACTION

Withdraw Money from Hossein's Saving Account

Deposit Money to Hossein's Checking Account

END TRANSACTION

Physical Level × File × ACID

77

Consistency

Respect constraints or expectations among data instances

- CONSTRAINT#1: All Movies Must Have ReleaseDate ≥ 1890
- CONSTRAINT#2: All Movies Must Have at Least One Director
- CONSTRAINT#3: All Movies Must Have UNIQUE Title

Physical Level × File × ACID

78

Consistency

Respect constraints or expectations among data instances

```
BEGIN TRANSACTION
```

```
  INSERT Movie
```

```
    INSERT Movie's distributor Company IF NOT EXIST
```

```
END TRANSACTION
```

Transaction Conflicts with CONSTRAINT#2.

ROLLBACK any changes.

Physical Level × File × ACID

79

Consistency

Respect constraints or expectations among data instances

```
BEGIN TRANSACTION
```

```
  INSERT Movie
```

```
    Error: INSERT Movie's Director IF NOT EXIST
```

```
    INSERT Movie's distributor Company IF NOT EXIST
```

```
END TRANSACTION
```

```
Transaction Committed.
```

```
2 records affected.
```

```
Partial Execution of Transaction → Transaction Conflicts with CONSTRAINT#2
```

Inconsistency was due to lack of atomicity.

ACID properties are not orthogonal (independent).

Physical Level × File × ACID

80

Isolation

Transaction appear to be executed as if no other transaction is executing at the same time

Multi-user environment!

Physical Level × File × ACID

81

Isolation

Transaction appear to be executed as if no other transaction is executing at the same time

Transaction Queue (Sequential Execution, One-at-a-time)		
T1	W1	T1.1. <i>INSERT 2001: A Space Odyssey</i>
	W2	T1.2. <i>INSERT Stanley Kubrick IF NOT EXIST</i>
	W3	T1.3. <i>INSERT MGM IF NOT EXIST</i>
T2	W4	T2.1. <i>INSERT A Clockwork Orange</i>
	W5	T2.2. <i>INSERT Stanley Kubrick IF NOT EXIST</i>
	W6	T2.3. <i>INSERT Warner Bros IF NOT EXIST</i>
T3	W7	T3.1. <i>UPDATE Director SET PlaceOfBirth = "United States of America" WHERE PlaceOfBirth = "USA"</i>
T4	W8	...

Physical Level × File × ACID

82

Isolation

Transaction appear to be executed as if no other transaction is executing at the same time

Transaction Queue (Parallel Execution)			
W1	T1.1. <i>INSERT 2001: A Space Odyssey</i>	T2.1. <i>INSERT A Clockwork Orange</i>	
W2	T1.2. <i>INSERT Stanley Kubrick IF NOT EXIST</i>	T1.3. <i>INSERT MGM IF NOT EXIST</i>	T2.3. <i>INSERT Warner Bros IF NOT EXIST</i>
W3	T2.2. <i>INSERT Stanley Kubrick IF NOT EXIST</i>		
W4	T3.1. <i>UPDATE Director SET PlaceOfBirth = "United States of America" WHERE PlaceOfBirth = "USA"</i>		

Physical Level × File × ACID

83

Isolation

Transaction appear to be executed as if no other transaction is executing at the same time

Transaction Queue (Parallel Execution)		
W1	T1.1. <i>INSERT 2001: A Space Odyssey</i>	T2.1. <i>INSERT A Clockwork Orange</i>
W2	T1.2. <i>INSERT Stanley Kubrick IF NOT EXIST</i>	T2.2. <i>INSERT Stanley Kubrick IF NOT EXIST</i>
W3	T1.3. <i>INSERT MGM IF NOT EXIST</i>	T2.3. <i>INSERT Warner Bros IF NOT EXIST</i>
W4	T3.1. <i>UPDATE Director SET PlaceOfBirth = "United States of America" WHERE PlaceOfBirth = "USA"</i>	

What's the Result? **Duplicate Stanley Kubrick!**

Physical Level × File × ACID

84

Durability

Once a transaction has committed, the effect must never be lost

Function Exception

Program Crash

System Crash: files might not be durable at this level.

"Apocalypse Now"



Starchild: Database Management System

Database Management Systems (DBMS) 86

A Software System

Provides General Purpose, Efficient, Convenient, & Safe Multi-user Storage of and Access to Massive Amounts of Persistent Data

ACID Properties Guaranteed!

Database Management Systems (DBMS) 87

General Purpose

Different Applications: Moviesion, Library,

Safe

From System Failures, From Malicious Users

Convenient

Simple DDL & DML

Efficient

Don't Search All Files in Order to Get a Record

ACID

Atomicity, Consistency, Isolation, Durability

Database Management Systems (DBMS) 88

	People (Role)
Is this course about building a DBMS?	No! DBMS Implementors
Is this course about maintaining a DBMS?	No! DB Administers (DBA)
Is this course about <u>using</u> a DBMS?	<u>Yes</u> . DB Designers & User

DBMS Products

89

Company	DBMS
Oracle	Oracle, <u>MySQL</u>
Microsoft	MySQLServer
IBM	DB2
OpenLink Software	Virtuoso Universal Server
Apache Software Foundation	CouchDB

