



# Today

1



Data Modeling  
in  
RDBMS

Real World Entity

Conceptual Level | Entity-Relationship Model (E/R)

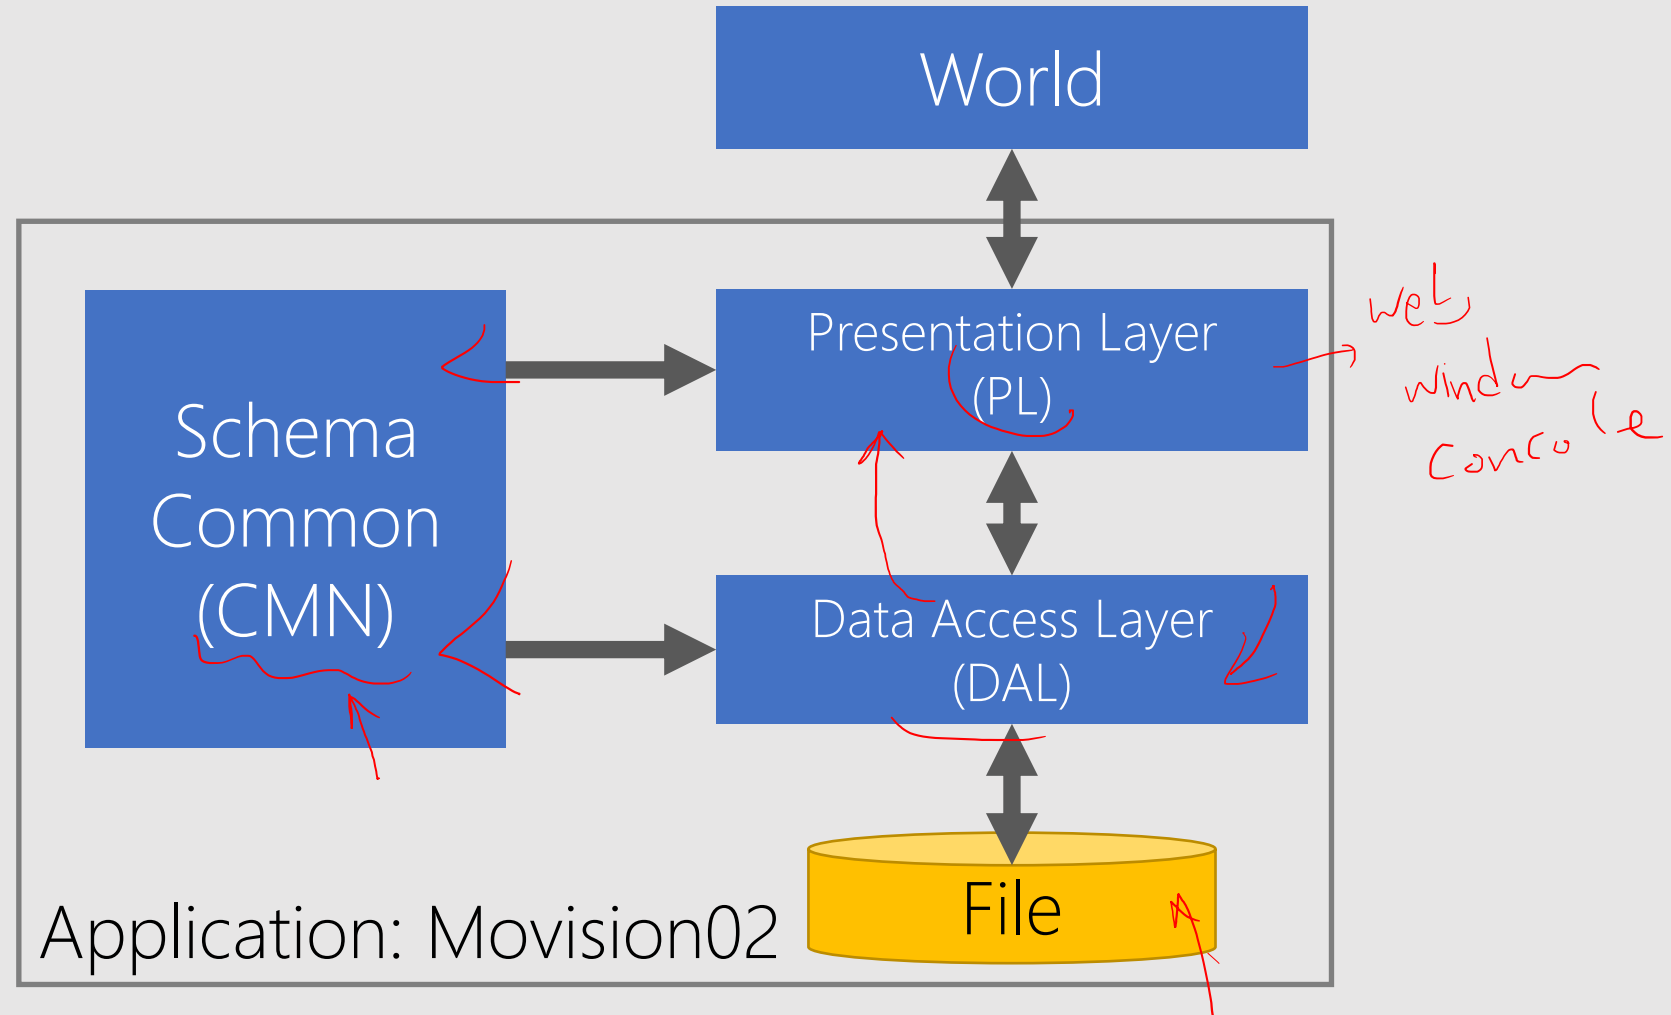
| Logical Level | Relational Model

| Physical Level | SQL

Computable Entity


# Physical Level × File

2



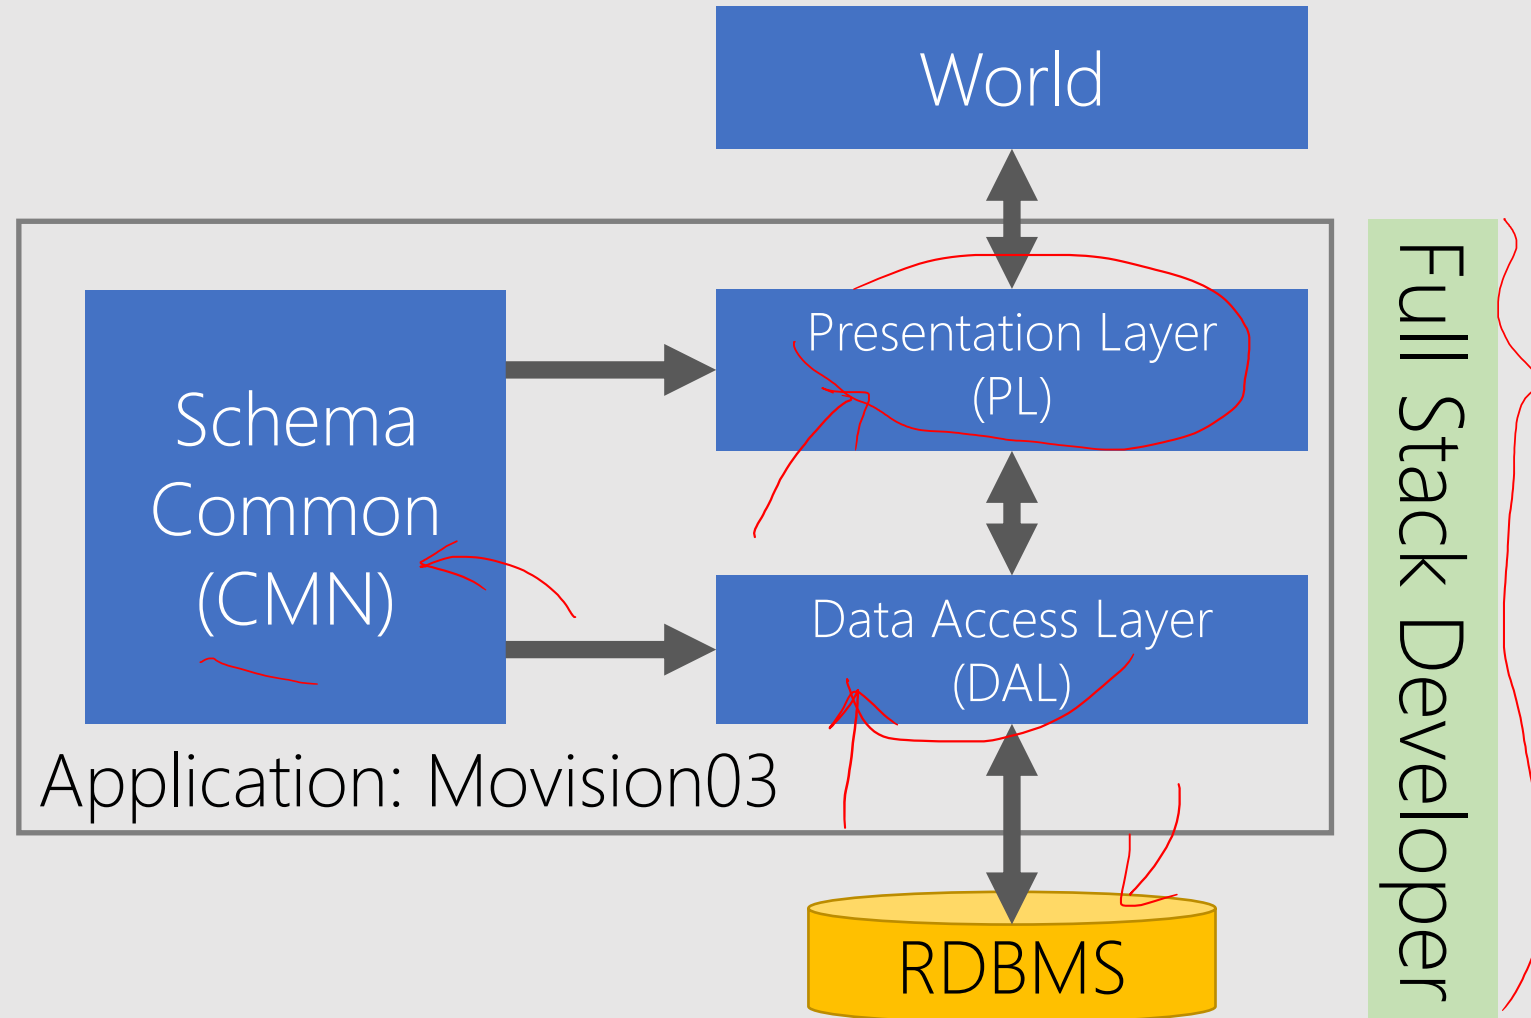
Physical Level × RDBMS

3



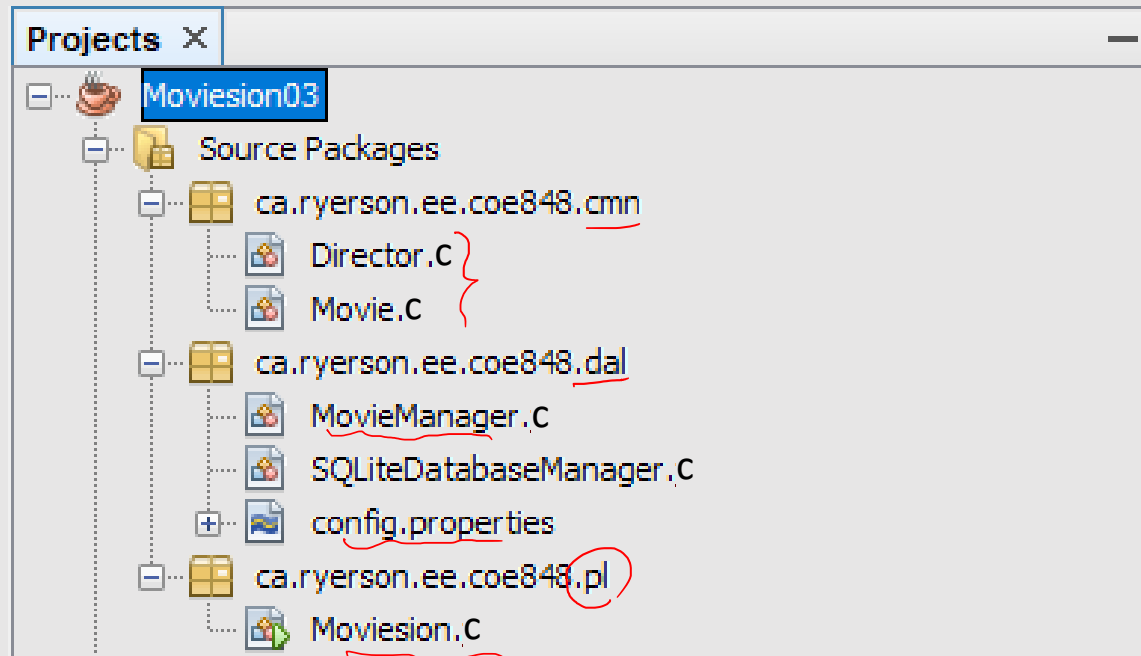
# Physical Level × RDBMS

4



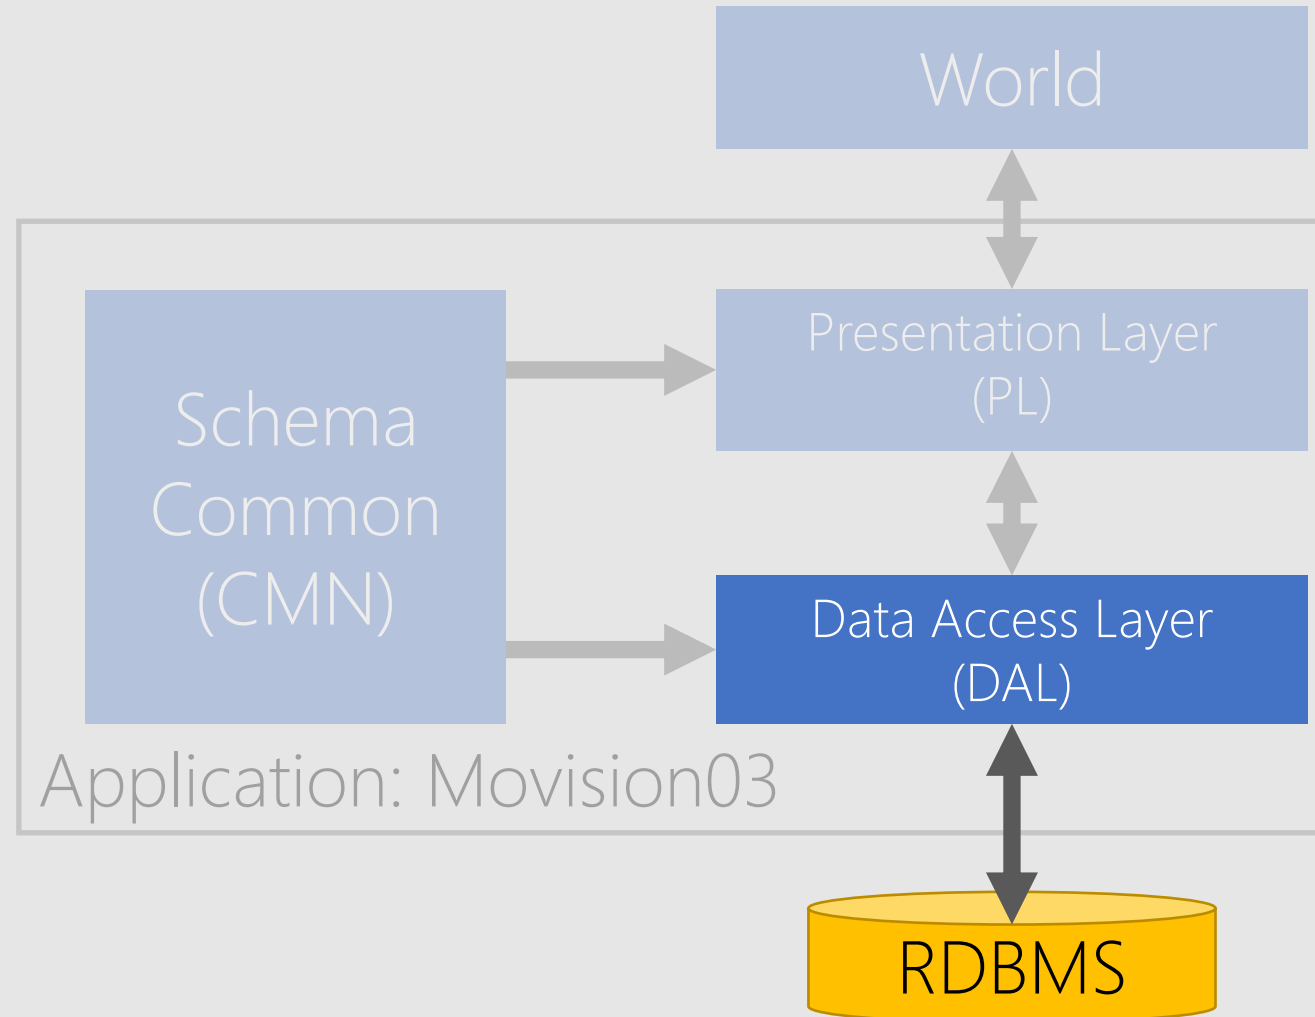
# Physical Level × RDBMS

5



# Physical Level × ORM

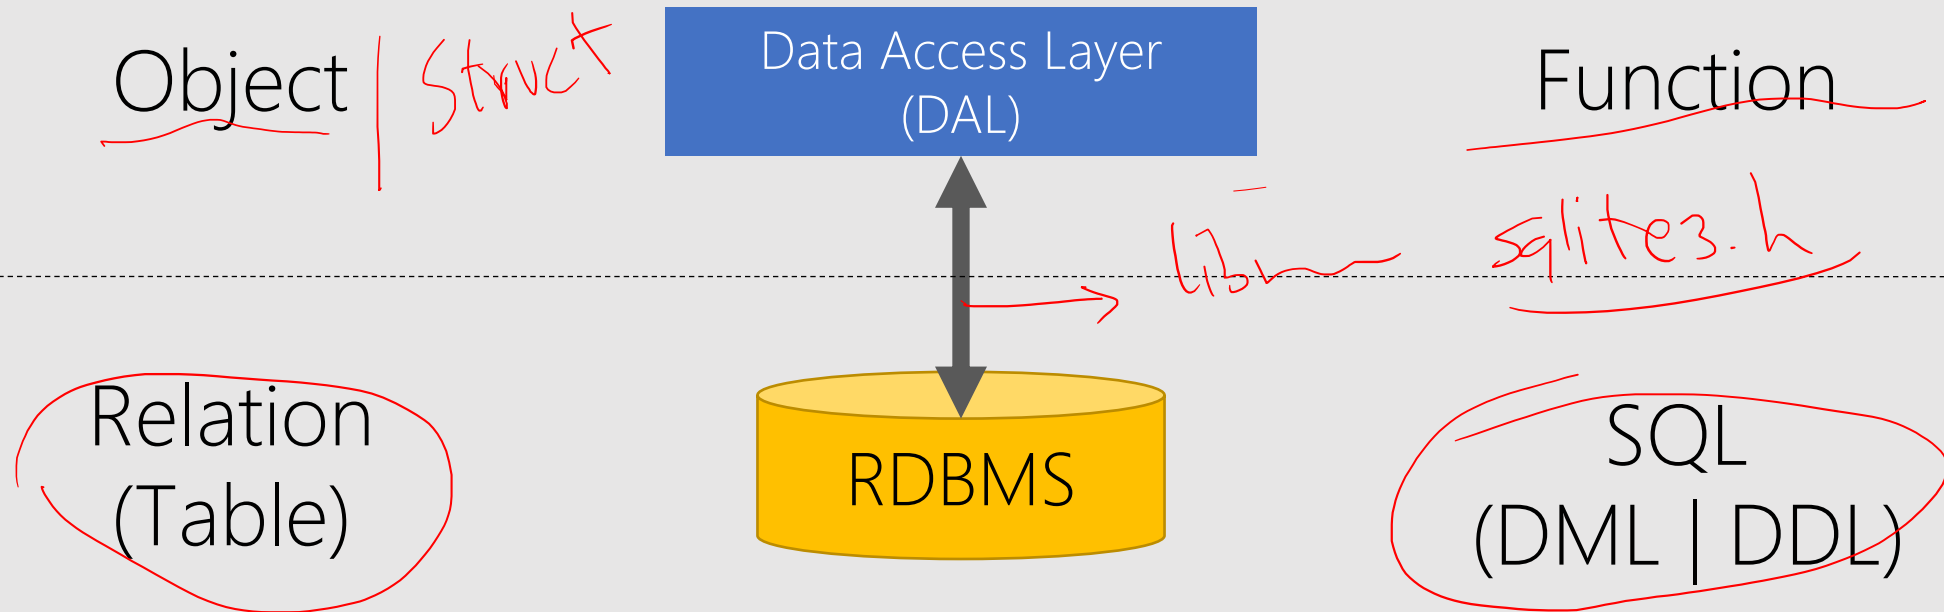
6



# Physical Level × ORM

7

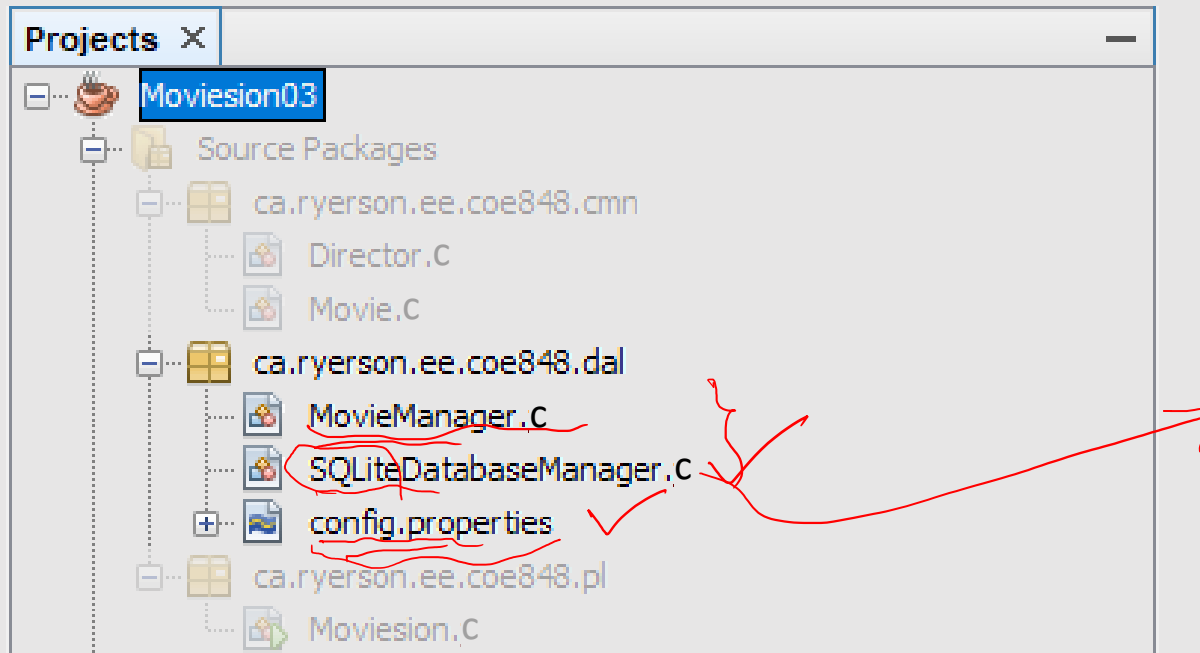
## Object Relational Mapping (ORM)






# Physical Level × ORM

8



# Physical Level × ORM

9




Settings to connect to DBMS:

DBMS Location  
DBMS Username  
DBMS Password  
DB Name

# Physical Level × ORM

10



Anything related to DBMS connection:

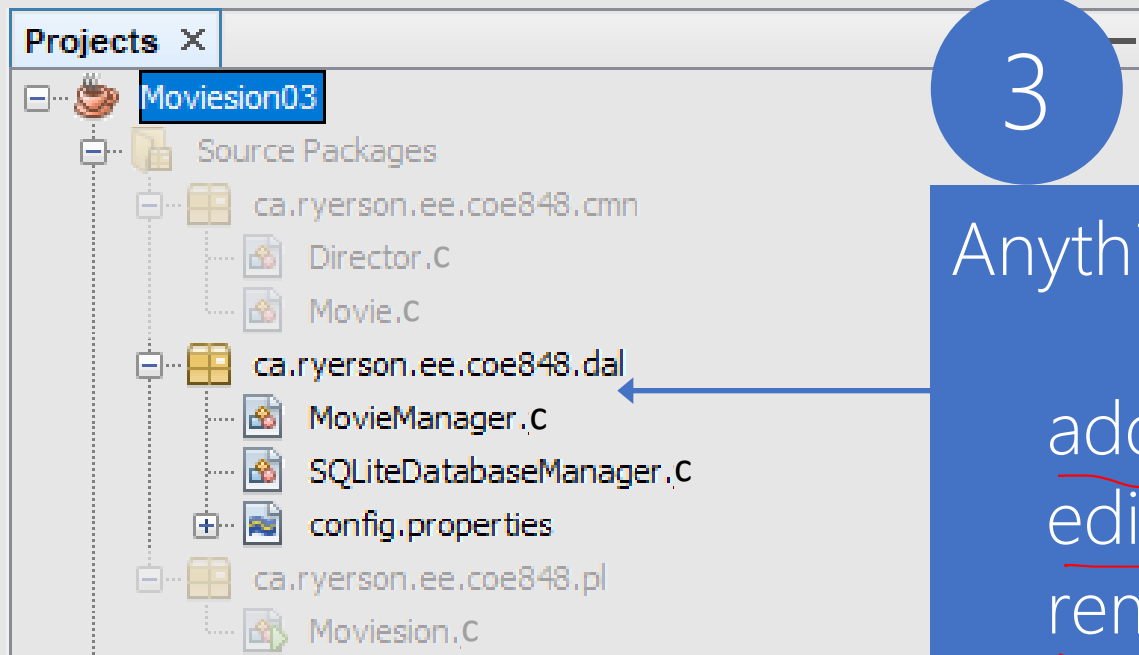
Loading SQLite library

Loading settings from config file

Create a connection

# Physical Level × ORM

11



3

Anything related to Movie (ORM):

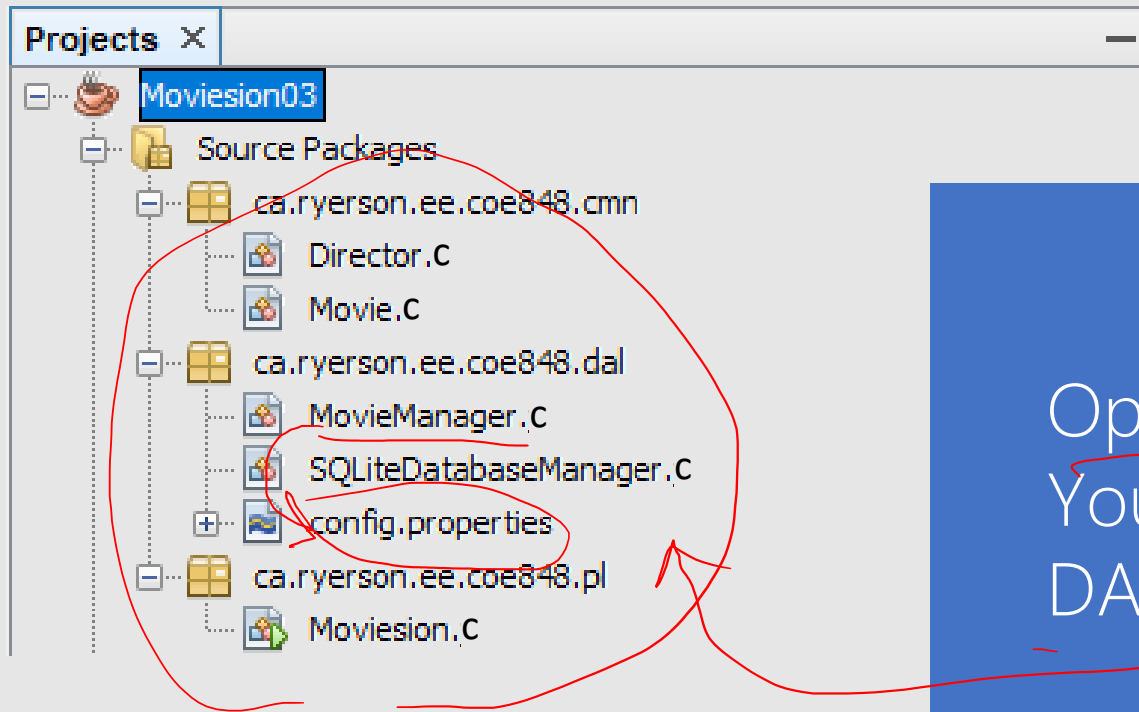
addMovie()  
editMovie()  
removeMovie()  
getMovie()

by Id (int)

INSERT  
UPDATE  
DELETE  
SELECT

# Physical Level × ORM

12




Optional:  
You can follow 3-layered architecture (PL,  
DAL, DB) for your Lab5 project.

# DB vs. APP Level Processing


13

How many movies do we have?



# DB Level Processing

14



How many movies do we have?


In SQL: what is the movie count?

- I) Count the movies
- II) Return a single number

*lo: # movie?*  
*CountMovie()*  
*return*

# APP Level Processing

15



How many movies do we have?

- I) Get all movies
- II) Count the movies

Return all movies

*for (i=0; i<n; i++)*  
*wh*



# DB vs. APP Level Processing

16

## APP Level

+:


Only simple SQL

Able to do very complex tasks

-:

Slow, moving all data to app. level

Waste of network bandwidth



## DB Level

+:

Fast, no need to move data

Fast, DBMS is a powerful machine

-:

Master SQL language

Not able to do very complex tasks

## *Ad hoc* SQL Query



17

Ad hoc query is created to obtain info as need arises,  
e.g., which director has made the most movies?

Contrast with a query that is predefined & routinely  
processed,

e.g., `INSERT`, `UPDATE`, `DELETE`, `SELECT` by Id



Eternal Sunshine of the Spotless Mind - Michel Gondry, 2004  
<https://www.youtube.com/watch?v=07-QBnEkgXU>

Indexing is finding the whole information quicker using only part of it.

Part of information is called Search Key.  
It points to the whole information.



Primary — Key  
Surrogate — Key  
Candidate — Key  
Foreign — Key

# Index

2

How to find a webpage in WWW?

Search Key = 'UWindsor'

Whole Information = <https://www.uwindsor.ca>

>  $10^{10}$  seconds by no index, traverse all webpages

< 0.31 seconds by Google

~ 0 seconds by ?

# SQL × INDEX

3

{  
SELECT \* FROM Director WHERE Id=1  
SELECT \* FROM Director WHERE LastName='Kubrick'  
SELECT \* FROM Director WHERE LastName='Kubrick' AND FirstName = 'Stanley'

# SQL × INDEX

4

CREATE [UNIQUE] INDEX IndexName ON TableName (c1, c2, ...);



• Could be any name, but by convention we follow this:

IX\_ColumnName1\_ColumnName2\_...

UIX\_ColumnName1\_ColumnName2\_...

UNIQUE INDEX does not allow duplicate in indexed columns.  
A way to create a candidate key set of columns in a table.

# SQL × INDEX

5

SELECT \* FROM Director WHERE Id = 1

CREATE UNIQUE INDEX UIX\_Id ON Director(Id)

$O(20 \text{ million})$

$O(1)$

By default, most DBMSs CREATE UNIQUE INDEX on primary key set of a table.



# SQL × INDEX

6

What other columns of a table should to be indexed?

- Those columns of table that appears a lot in WHERE clause.
- The search key of the table to find a single or range of rows.

It's a tuning task: → DBA



- After the DB goes under heavy load DB designer need to increase retrieval speed.
- Recently is done automatically by DBMS

# SQL × INDEX

 `SELECT * FROM Director WHERE LastName = 'Kubrick'`  
`CREATE INDEX IX_LastName ON Director(LastName)`  


P2

7

 8. Fun man  
by Lan;  


# SQL × INDEX

8

ALTER TABLE TableName ADD [UNIQUE] INDEX IndexName ON (c1, c2, ...);

ALTER TABLE TableName DROP INDEX IndexName;



# SQL × INDEX

9

SELECT \* FROM Director WHERE LastName = 'Kubrick' AND FirstName = 'Stanley'

Which one?

- A) CREATE INDEX IX\_LastName\_FirstName ON Director(LastName, FirstName)
- B) CREATE INDEX IX\_FirstName\_LastName ON Director(FirstName, LastName)
- C) CREATE INDEX IX\_FirstName ON Director(FirstName)
- D) CREATE INDEX IX\_LastName ON Director(LastName)
- E) All
- F) A & B are the same

# DBMS × INDEX

10

As far as DB designer is concerned, knowing how to **CREATE** | **ADD** | **DROP INDEX** in SQL is more than enough.

However, knowing the implementation details inside DBMS helps DB designer with right decisions about indexing.

# DBMS × INDEX × Binary Search Tree

11

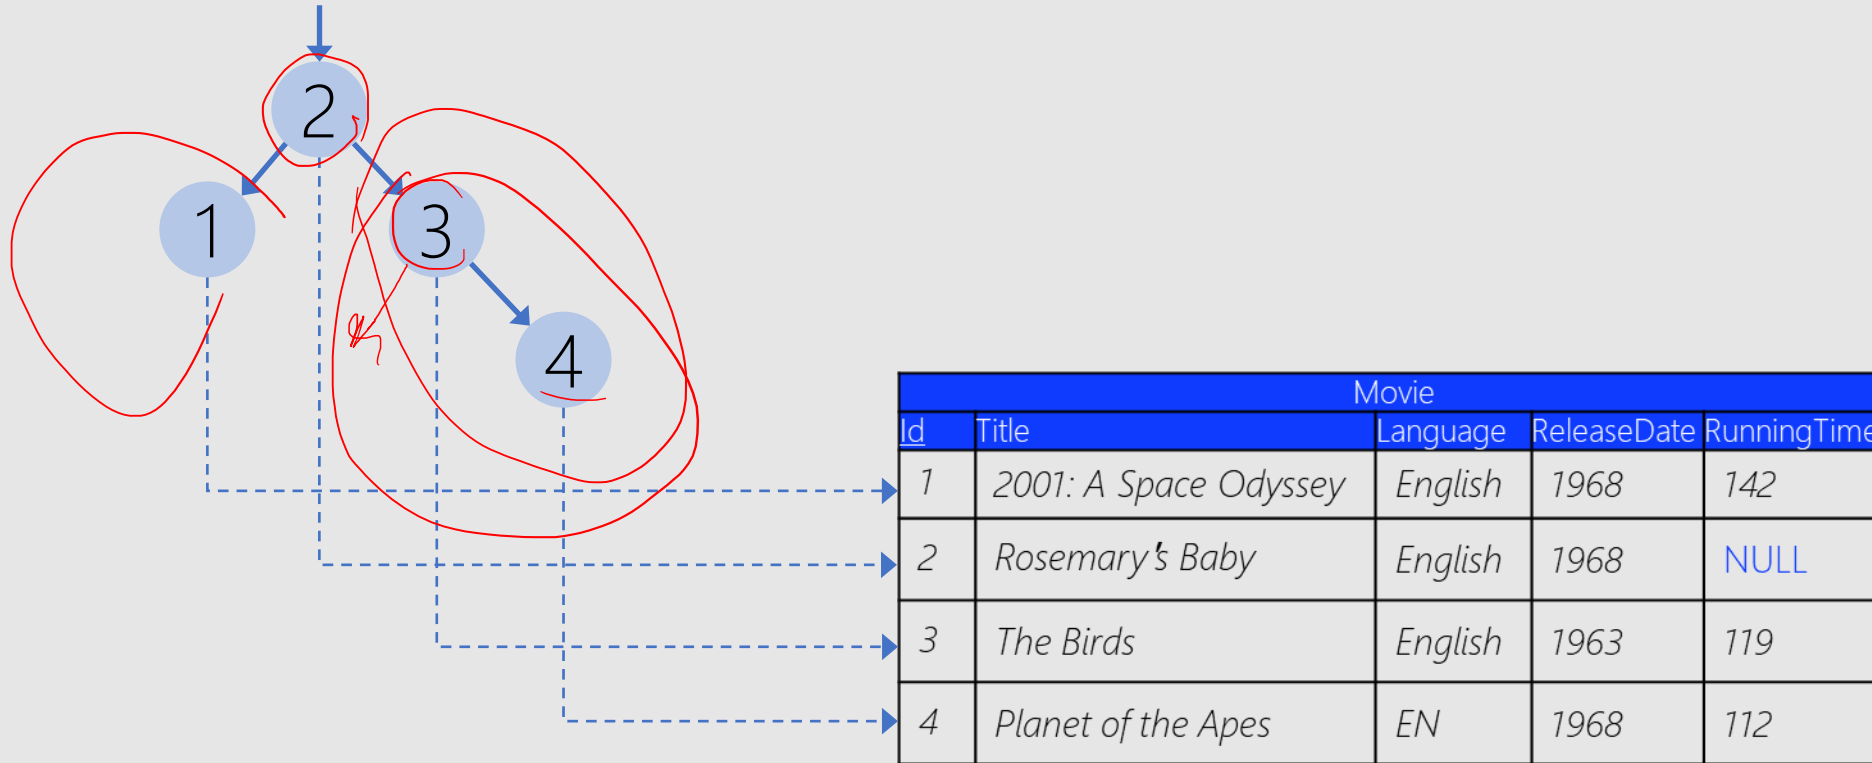
Movie				
Id	Title	Language	ReleaseDate	RunningTime
1	2001: A Space Odyssey	English	1968	142
2	Rosemary's Baby	English	1968	NULL
3	The Birds	English	1963	119
4	Planet of the Apes	EN	1968	112

SELECT \* FROM Movie WHERE Id = 1

- ✓ A. Sequential search, check all movies' Id with the given Id, i.e., 1
- ⓑ. Binary search, after sorting elements in the list by Id →  $\log_2 n$
- Ⓒ. Having an index structure:
  - Ⓐ. Creating a Binary Search Tree (BST)


# DBMS × INDEX × Binary Search Tree

12



# DBMS × INDEX × Binary Search Tree

13



Average

Search  $O(\log n)$

Insert  $O(\log n)$

Delete  $O(\log n)$

Worst (When?)

Search  $O(n)$

Insert  $O(n)$


Delete  $O(n)$

Movie				
Id	Title	Language	ReleaseDate	RunningTime
1	2001: A Space Odyssey	English	1968	142
2	Rosemary's Baby	English	1968	NULL
3	The Birds	English	1963	119
4	Planet of the Apes	EN	1968	112



# DBMS × INDEX × Binary Search Tree

14



Average

Search  $O(\log n)$

Insert  $O(\log n)$

Delete  $O(\log n)$

Worst (When?)

Search  $O(n)$

Insert  $O(n)$


Delete  $O(n)$

Movie				
Id	Title	Language	ReleaseDate	RunningTime
1	2001: A Space Odyssey	English	1968	142
2	Rosemary's Baby	English	1968	NULL
3	The Birds	English	1963	119
4	Planet of the Apes	EN	1968	112

Overhead:

Each DML on the table needs additional DML on indexes of the table by DBMS

# DBMS × INDEX × Balanced Binary Tree 15



Average

Search  $O(\log n)$

Insert  $O(\log n)$

Delete  $O(\log n)$

Worst

Search  $O(\log n)$

Insert  $O(\log n)$

Delete  $O(\log n)$

Movie				
Id	Title	Language	ReleaseDate	RunningTime
1	2001: A Space Odyssey	English	1968	142
2	Rosemary's Baby	English	1968	NULL
3	The Birds	English	1963	119
4	Planet of the Apes	EN	1968	112

Overhead:

Each DML on the table needs additional DML on indexes of the table by DBMS

DBMS × INDEX × B-tree

16

Balanced Multi-way Tree

Bayer, R.; McCreight, E. (1972)

*Organization and Maintenance of Large Ordered Indexes*  
Acta Informatica, 1 (3): 173–189

@Boeing




# DBMS × INDEX × B-tree

17

B-tree of order  $m$  (branching factor) is a tree in which:

- $0 \leq \# \text{keys in root} \leq (m-1)$   $\rightarrow m$
- $0 \leq \# \text{subtrees in root} \leq m$
- $(\frac{1}{2} m) \leq \# \text{keys in other nodes} \leq (m-1)$
- $1 + (\frac{1}{2} m) \leq \# \text{subtrees in other nodes} \leq m$

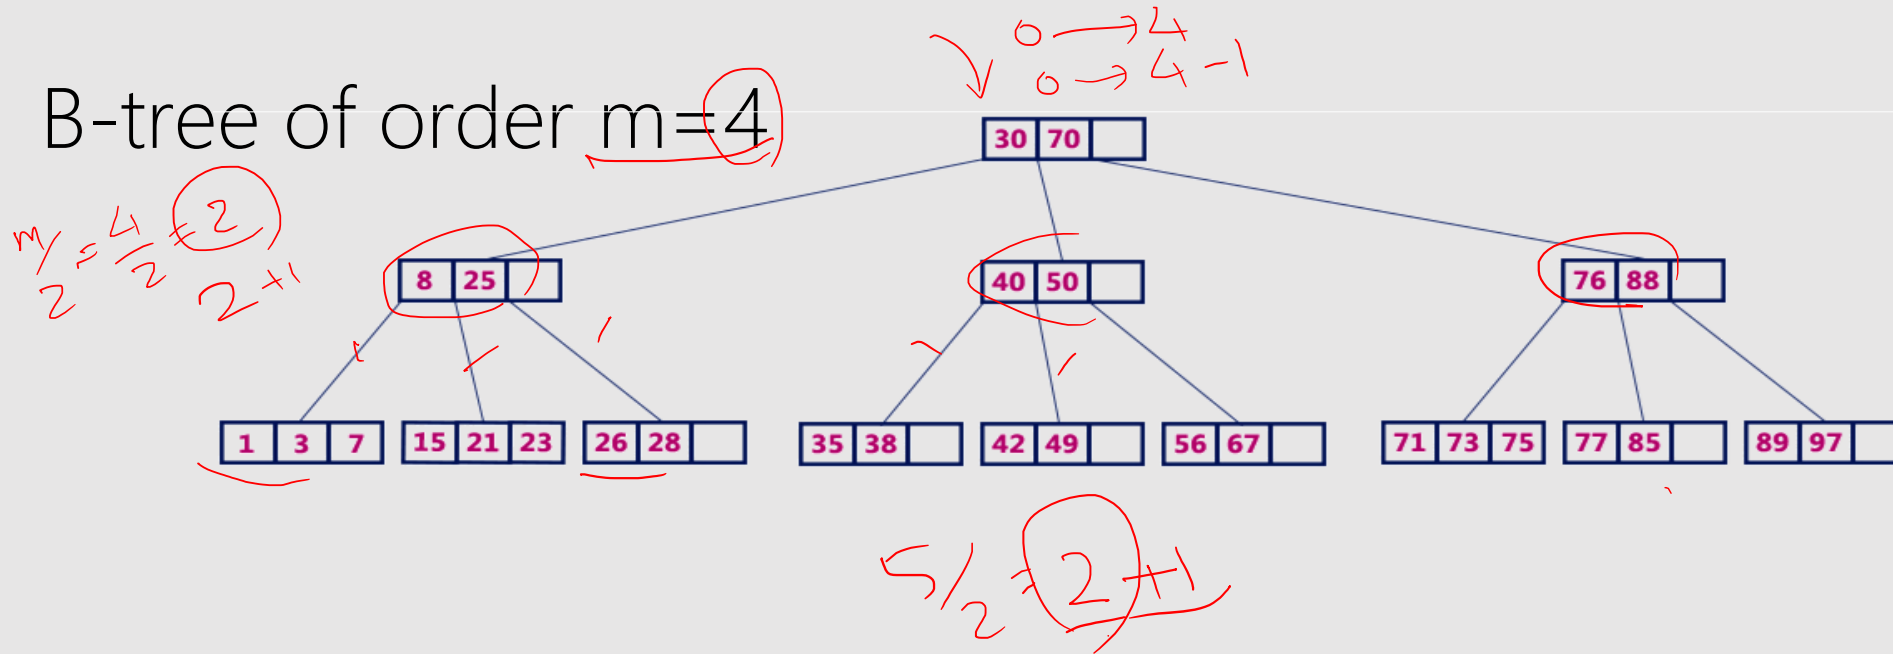
- The keys in each node are sorted.
- It is **balanced**!




# DBMS × INDEX × B-tree

18

B-tree of order  $m=4$



B-tree of order  $m=5$



# DBMS × INDEX × B-tree

19

The height  $h$  of a B-tree of order  $m$ , with a total of  $n$  keys:

$$\log_m^{(n+1)} \leq h \leq 1 + \log_{\lceil m/2 \rceil}^{(n+1/2)}$$

156

If  $m=300$  and  $n=16,000,000$  then  $h \approx 4$ .

i.e., the worst case finding a key in such B-tree requires ? accesses.