

"INSERT INTO T
VALUES (0, 2)"
✓ (@, 2)

Eternal Sunshine of the Spotless Mind - Michel Gondry, 2004
<https://www.youtube.com/watch?v=07-QBnEkgXU>

Indexing is finding the whole information quicker using only part of it.

Part of information is called Search Key.
It points to the whole information.



Primary — Key
Surrogate — Key
Candidate — Key
Foreign — Key

Index

2

How to find a webpage in WWW?

Search Key = 'UWindsor'

Whole Information = <https://www.uwindsor.ca>

> 10^{10} seconds by no index, traverse all webpages

< 0.31 seconds by Google

~ 0 seconds by ?

SQL × INDEX

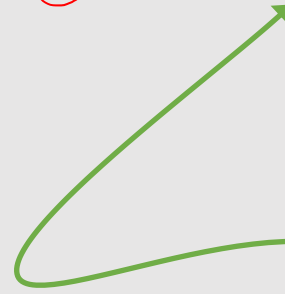
3

{
SELECT * FROM Director WHERE Id=1
SELECT * FROM Director WHERE LastName='Kubrick'
SELECT * FROM Director WHERE LastName='Kubrick' AND FirstName = 'Stanley'
}

SQL × INDEX

4

CREATE [UNIQUE] INDEX IndexName ON TableName (c1, c2, ...);



• Could be any name, but by convention we follow this:

IX_ColumnName1_ColumnName2_...

UIX_ColumnName1_ColumnName2_...

UNIQUE INDEX does not allow duplicate in indexed columns.

A way to create a candidate key set of columns in a table.

SQL × INDEX

5

SELECT * FROM Director WHERE Id = 1 $\rightarrow O(20 \text{ rows})$
CREATE UNIQUE INDEX UIX_Id ON Director(Id) $\leftarrow O(1)$

By default, most DBMSs CREATE UNIQUE INDEX on primary key set of a table.

SQL × INDEX

6

What other columns of a table should to be indexed?

- Those columns of table that appears a lot in WHERE clause.
- The search key of the table to find a single or range of rows.

It's a tuning task: → DBA

- After the DB goes under heavy load DB designer need to increase retrieval speed.
- Recently is done automatically by DBMS

SQL × INDEX

 `SELECT * FROM Director WHERE LastName = 'Kubrick'`
`CREATE INDEX IX_LastName ON Director(LastName)`


P2

7

⑧: Fun man
by Lan;

SQL × INDEX

8

ALTER TABLE TableName ADD [UNIQUE] INDEX IndexName ON (c1, c2, ...);

ALTER TABLE TableName DROP INDEX IndexName;



SQL × INDEX

9

SELECT * FROM Director WHERE LastName = 'Kubrick' AND FirstName = 'Stanley'

Which one?

- A) CREATE INDEX IX_LastName_FirstName ON Director(LastName, FirstName)
- B) CREATE INDEX IX_FirstName_LastName ON Director(FirstName, LastName)
- C) CREATE INDEX IX_FirstName ON Director(FirstName)
- D) CREATE INDEX IX_LastName ON Director(LastName)
- E) All
- F) A & B are the same

DBMS × INDEX

10

As far as DB designer is concerned, knowing how to **CREATE** | **ADD** | **DROP INDEX** in SQL is more than enough.

However, knowing the implementation details inside DBMS helps DB designer with right decisions about indexing.

DBMS × INDEX × Binary Search Tree

11

Movie				
Id	Title	Language	ReleaseDate	RunningTime
1	2001: A Space Odyssey	English	1968	142
2	Rosemary's Baby	English	1968	NULL
3	The Birds	English	1963	119
4	Planet of the Apes	EN	1968	112

SELECT * FROM Movie WHERE Id = 1

- ✓ A. Sequential search, check all movies' Id with the given Id, i.e., 1
- ⓑ. Binary search, after sorting elements in the list by Id → $\log_2 n$
- Ⓒ. Having an index structure:
 - Ⓐ. Creating a Binary Search Tree (BST)

DBMS × INDEX × Binary Search Tree

14

Average

Search $O(\log n)$

Insert $O(\log n)$

Delete $O(\log n)$

Worst (When?)

Search $O(n)$

Insert $O(n)$

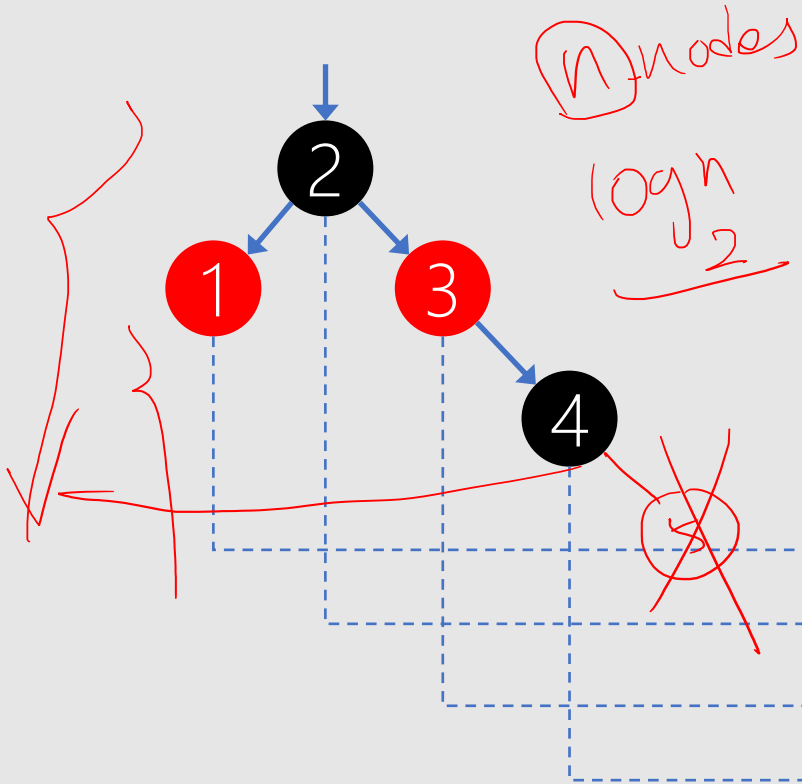
Delete $O(n)$

Movie				
Id	Title	Language	ReleaseDate	RunningTime
1	2001: A Space Odyssey	English	1968	142
2	Rosemary's Baby	English	1968	NULL
3	The Birds	English	1963	119
4	Planet of the Apes	EN	1968	112

Overhead:

Each DML on the table needs additional DML on indexes of the table by DBMS

DBMS × INDEX × Balanced Binary Tree 15



Average

Search $O(\log n)$

Insert $O(\log n)$

Delete $O(\log n)$

Worst

Search $O(\log n)$

Insert $O(\log n)$

Delete $O(\log n)$

Movie				
Id	Title	Language	ReleaseDate	RunningTime
1	2001: A Space Odyssey	English	1968	142
2	Rosemary's Baby	English	1968	NULL
3	The Birds	English	1963	119
4	Planet of the Apes	EN	1968	112

Overhead:

Each DML on the table needs additional DML on indexes of the table by DBMS

DBMS × INDEX × B-tree

16

Balanced Multi-way Tree



Bayer, R.; McCreight, E. (1972)

Organization and Maintenance of Large Ordered Indexes

Acta Informatica, 1 (3): 173–189

@Boeing



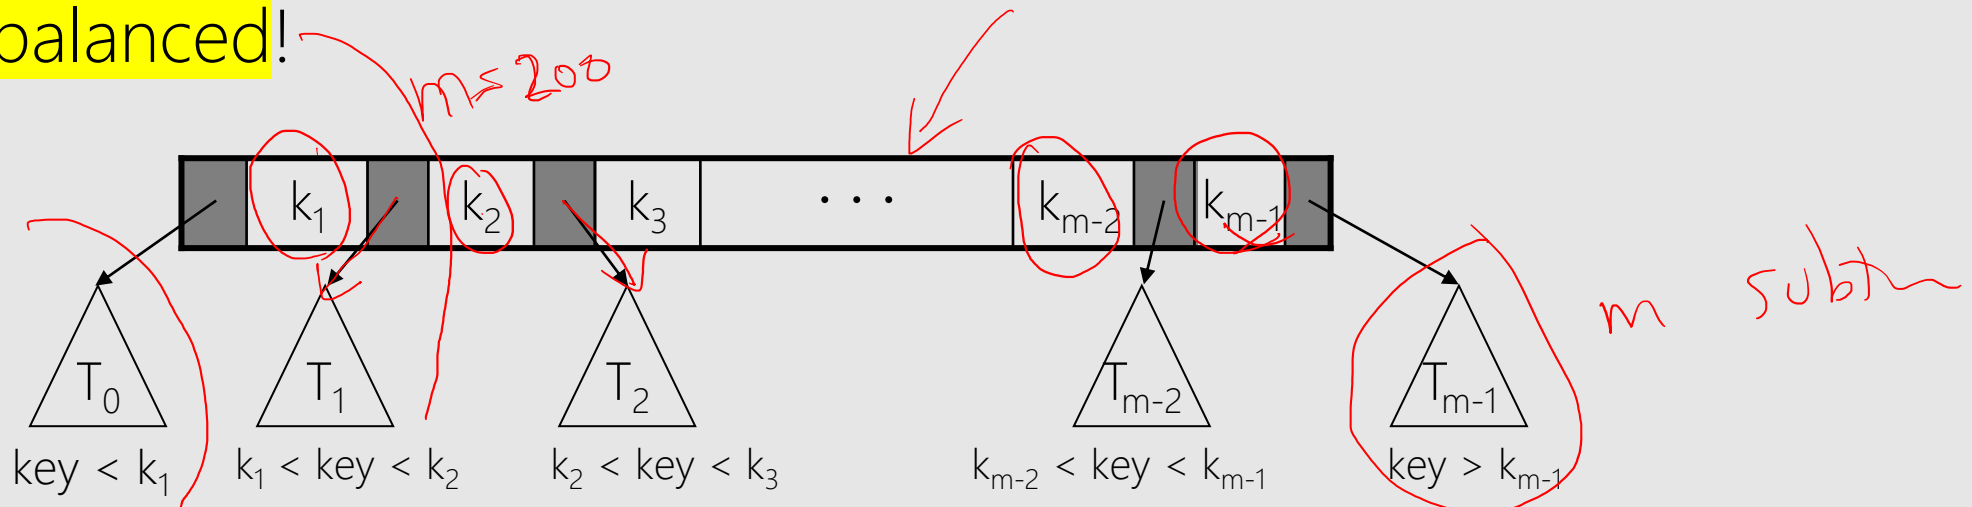
DBMS × INDEX × B-tree

17

B-tree of order m (branching factor) is a tree in which:

- $0 \leq \# \text{keys in root} \leq (m-1)$
- $0 \leq \# \text{subtrees in root} \leq m$
- $(\frac{1}{2} m) \leq \# \text{keys in other nodes} \leq (m-1)$
- $1 + (\frac{1}{2} m) \leq \# \text{subtrees in other nodes} \leq m$

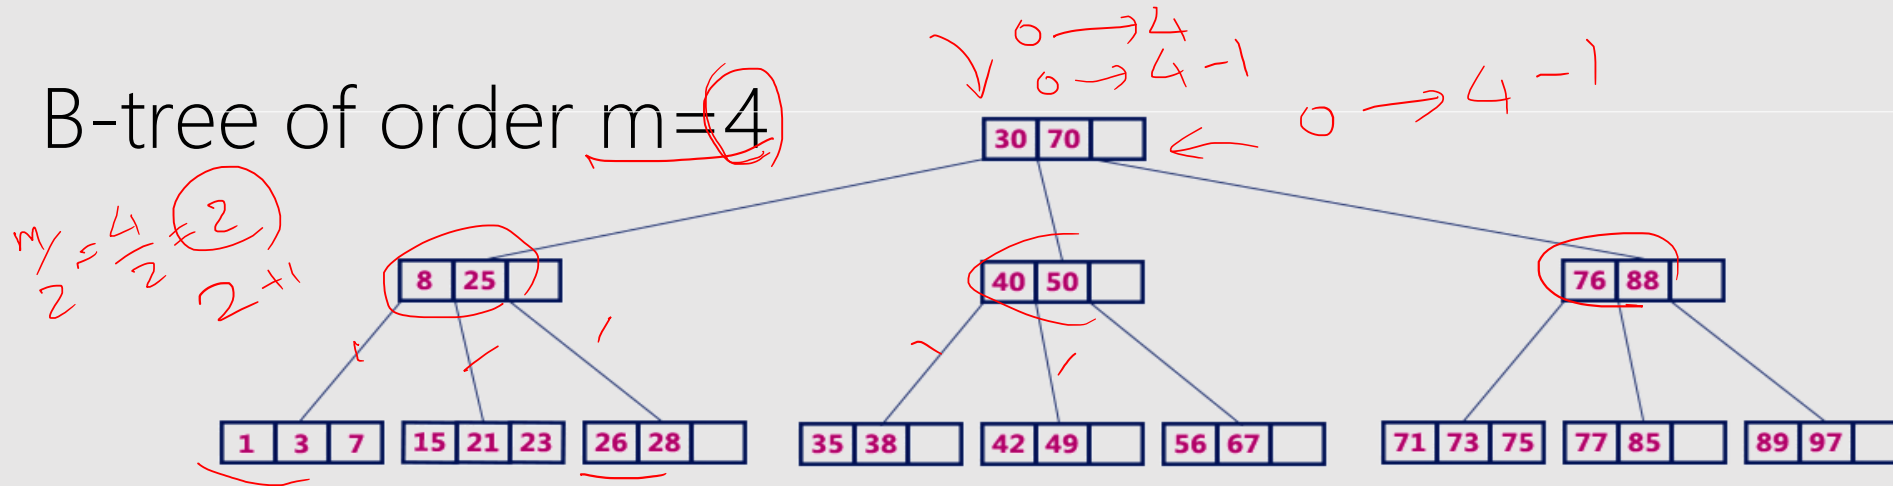
- The keys in each node are sorted.
- It is **balanced**!



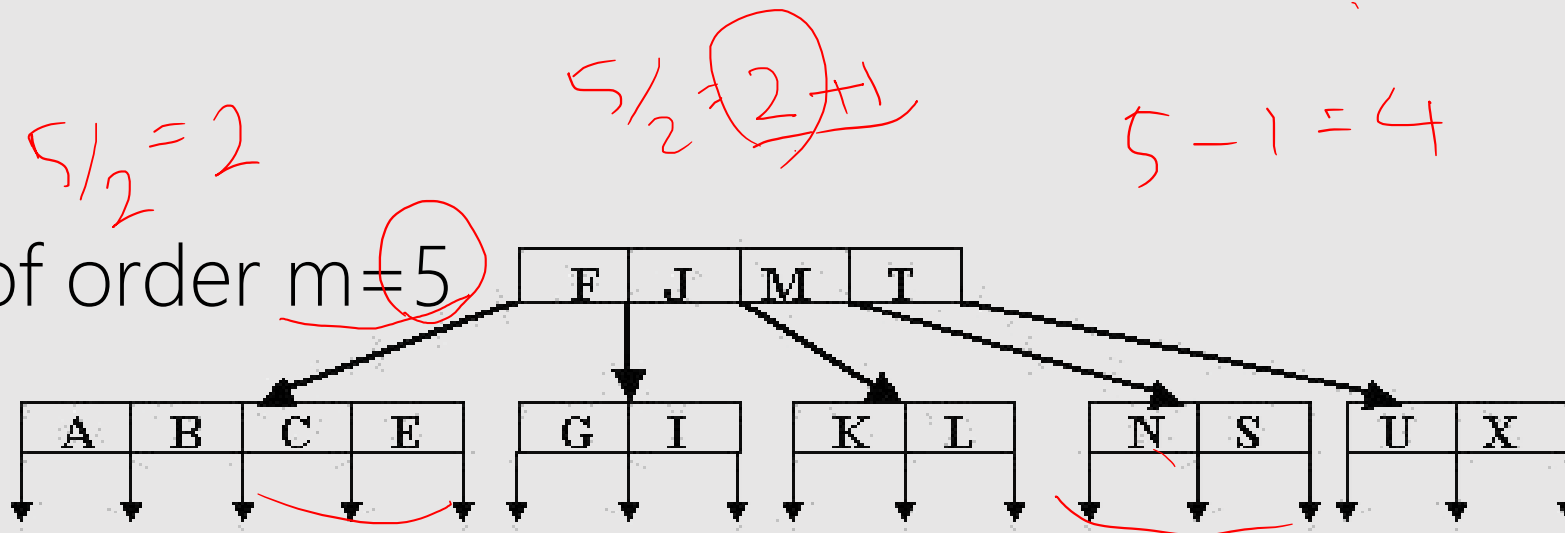
DBMS × INDEX × B-tree

18

B-tree of order $m=4$



B-tree of order $m=5$



DBMS × INDEX × B-tree

19

The height h of a B-tree of order m , with a total of n keys:

$$\log_m^{(n+1)} \leq h \leq 1 + \log_{\lceil m/2 \rceil}^{(n+1/2)}$$



If $m=300$ and $n=16,000,000$ then $h \approx 4$.

i.e., the worst case finding a key in such B-tree requires ? accesses.

$$\log_{300} 16,000,000 \sim 4$$

B-tree × DML

20

INSERT

DELETE

Overflow, more than m-1 keys

Underflow, less than m/2 keys

B-tree × INSERT

Empty B-tree of order $m=3$

Insert the keys 78, 52, 81, 40, 33, 90, 85, 20, and 38 in order



B-tree × INSERT

22

Empty B-tree of order $m=3$

Insert the keys 78, 52, 81, 40, 33, 90, 85, 20, and 38 in order

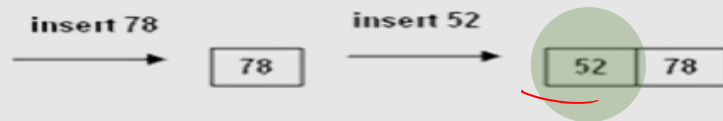


B-tree × INSERT

23

Empty B-tree of order $m=3$

Insert the keys 78, 52, 81, 40, 33, 90, 85, 20, and 38 in order



B-tree × INSERT

24

Empty B-tree of order $m=3$

Insert the keys 78, 52, 81, 40, 33, 90, 85, 20, and 38 in order

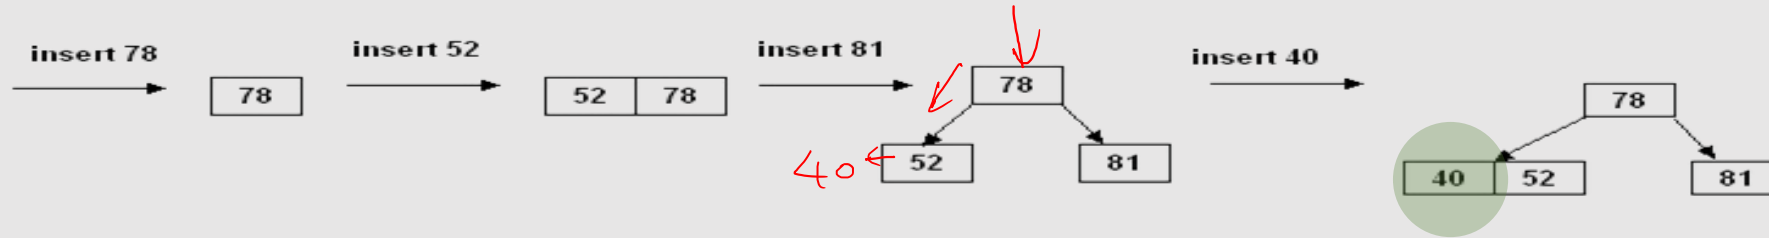


B-tree × INSERT

25

Empty B-tree of order $m=3$

Insert the keys 78, 52, 81, 40, 33, 90, 85, 20, and 38 in order

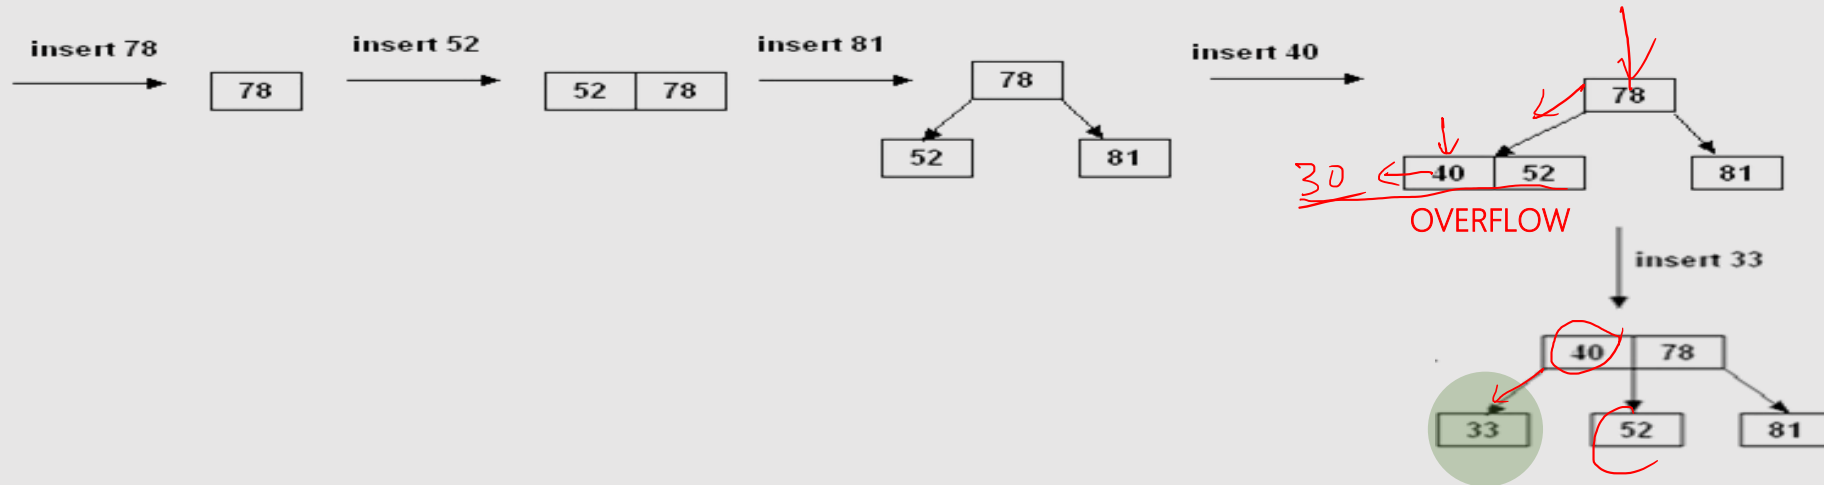


B-tree × INSERT

26

Empty B-tree of order $m=3$

Insert the keys 78, 52, 81, 40, 33, 90, 85, 20, and 38 in order

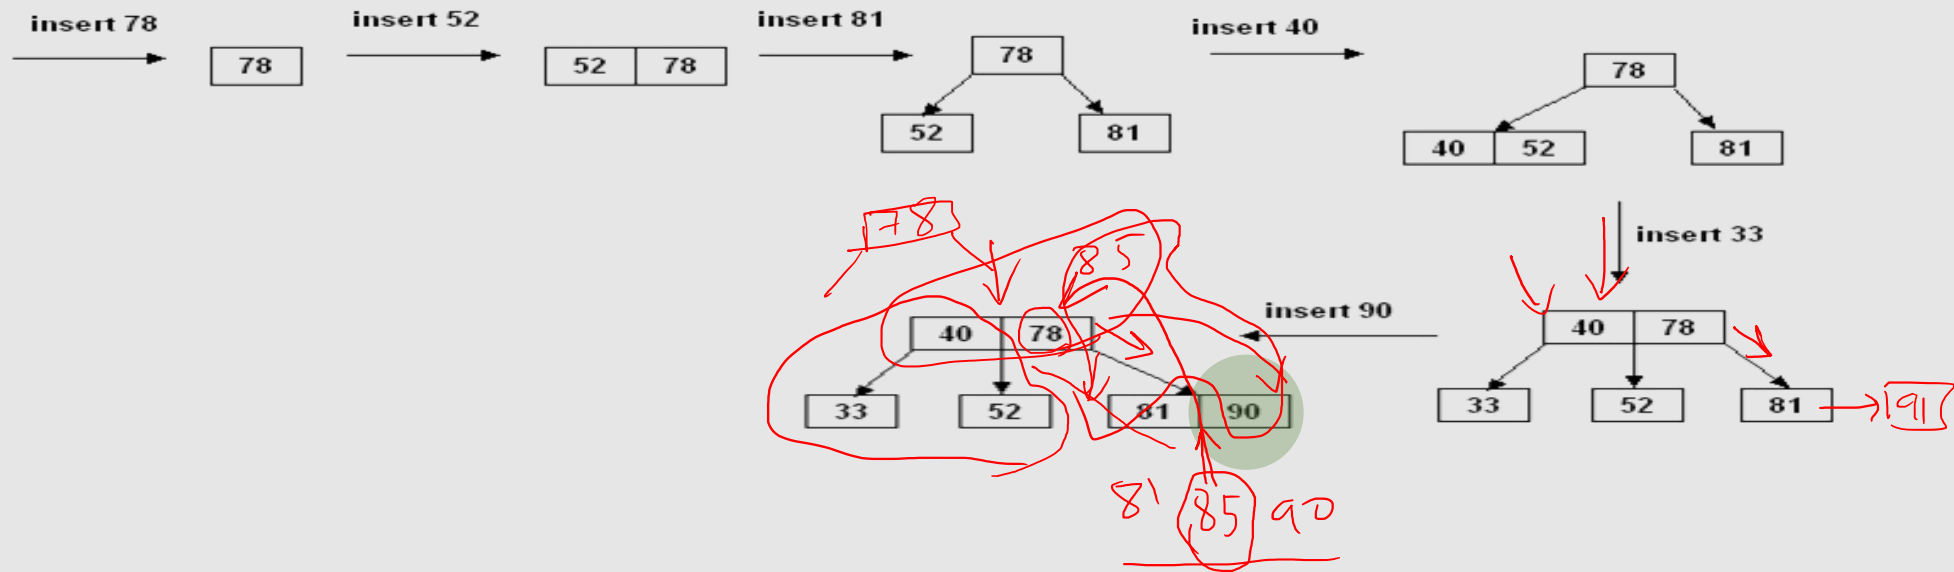


B-tree × INSERT

27

Empty B-tree of order $m=3$

Insert the keys 78, 52, 81, 40, 33, 90, 85, 20, and 38 in order

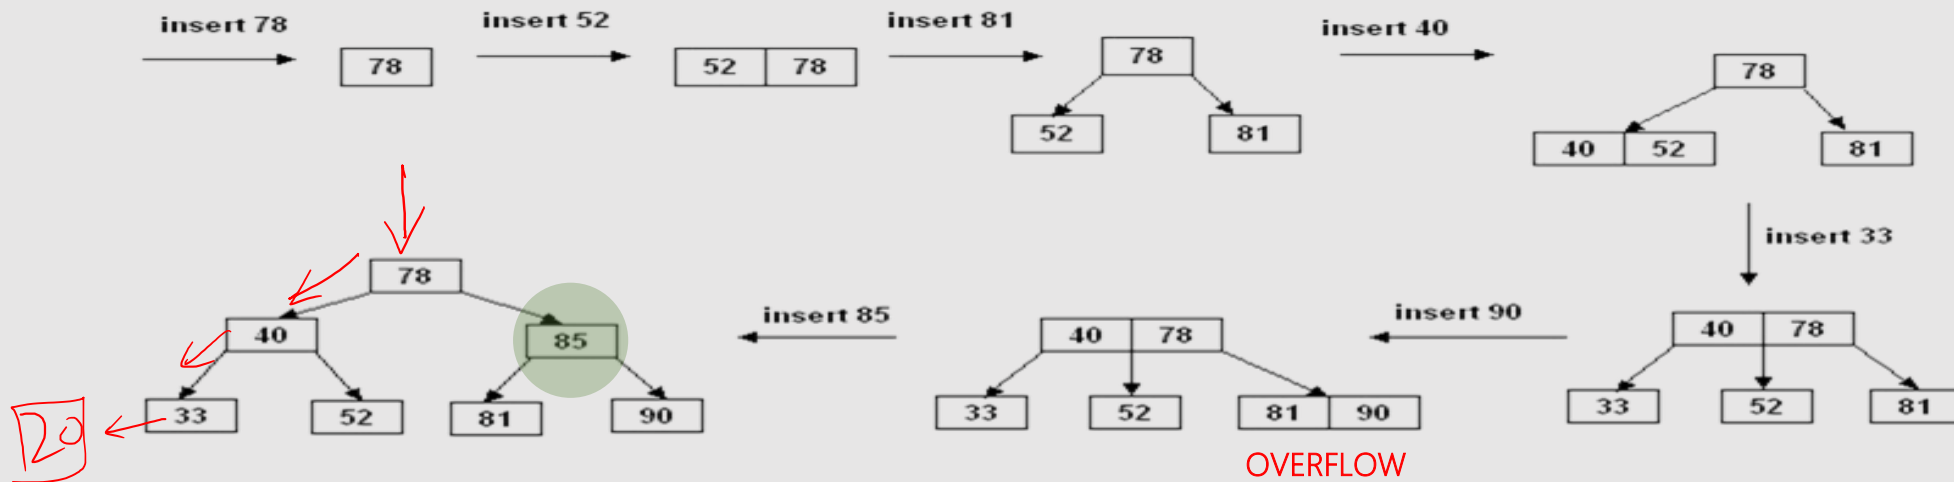


B-tree × INSERT

28

Empty B-tree of order $m=3$

Insert the keys 78, 52, 81, 40, 33, 90, 85, 20, and 38 in order

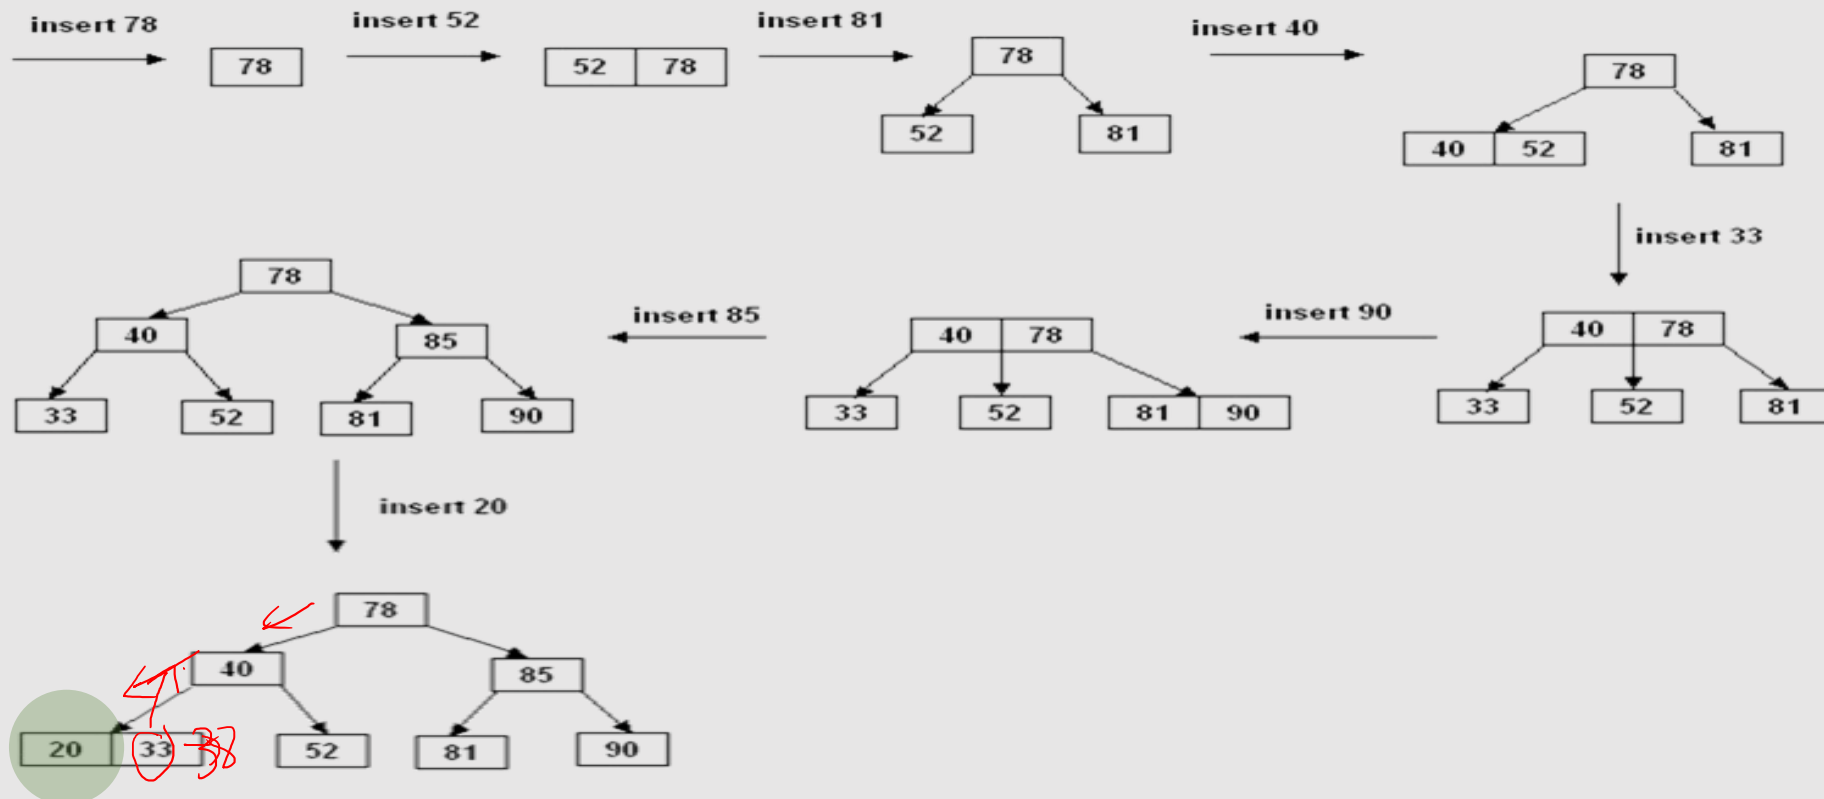


B-tree × INSERT

29

Empty B-tree of order $m=3$

Insert the keys 78, 52, 81, 40, 33, 90, 85, 20, and 38 in order

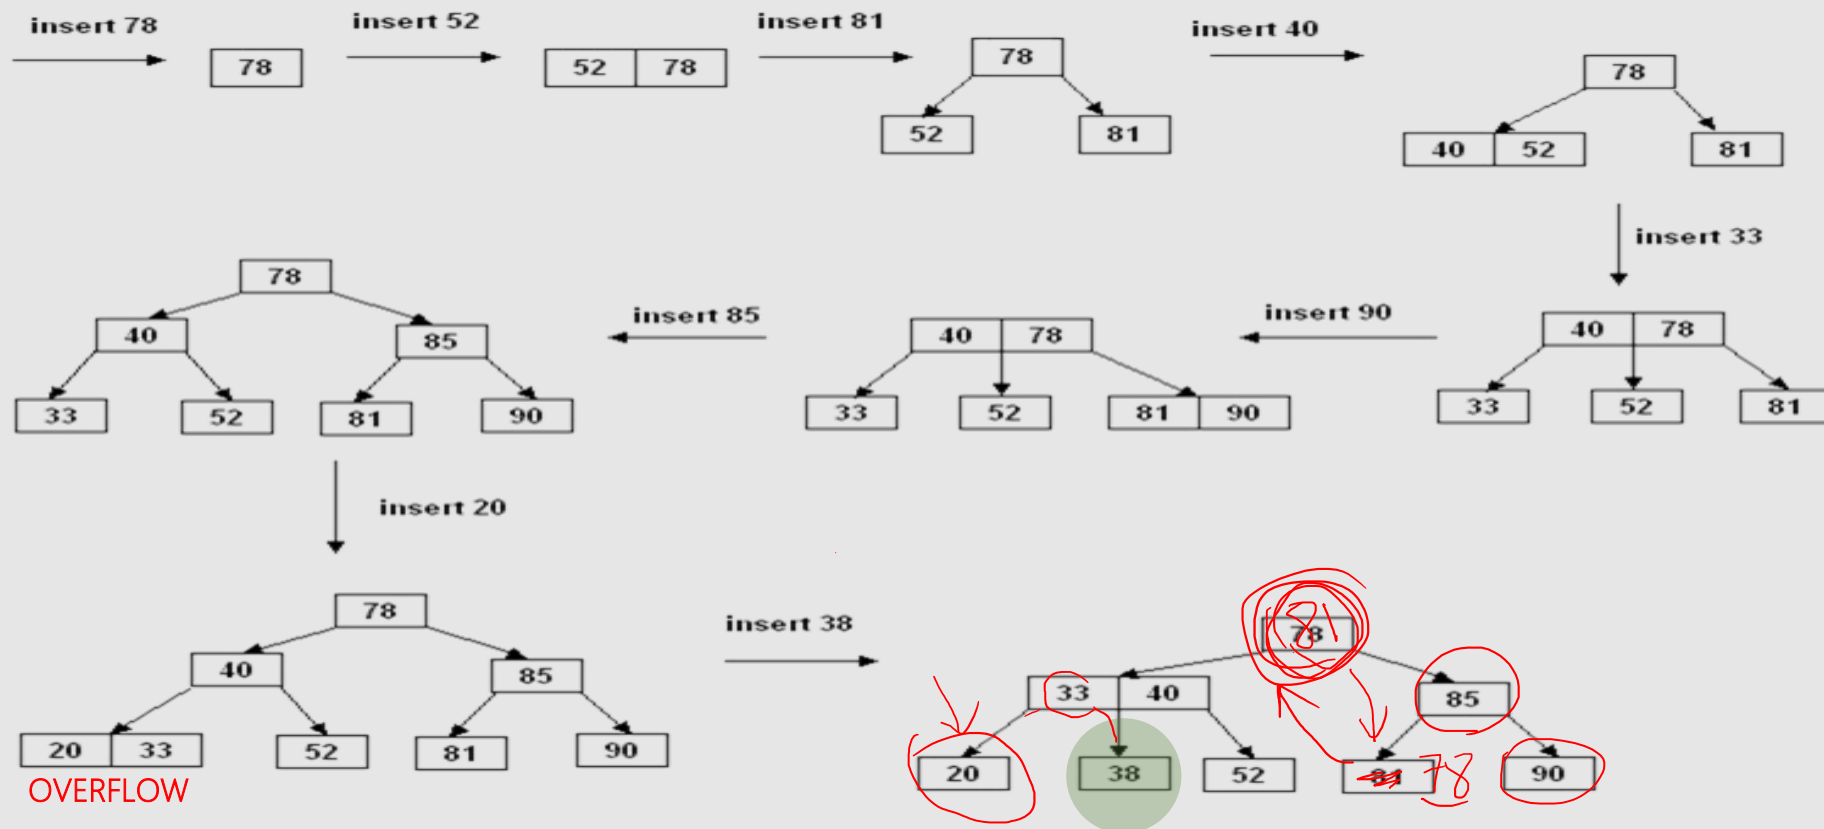


B-tree × INSERT

30

Empty B-tree of order $m=3$ $\frac{1}{2} = 1$

Insert the keys 78, 52, 81, 40, 33, 90, 85, 20, and 38 in order



B-tree × INSERT

31

INSERT(key):

Find the correct spot in a leaf node and do insert.

WHILE OVERFLOW

Split it into two

IF the node has parent

{ Insert the middle key to the parent (propagate)

Create right and left siblings

ELSE

Create a new root node

B-tree × DELETE

32

DELETE can happen at leaf or non-leaf!

IF the key is in a leaf node

Delete the key ✗

ELSE //non-leaf

} Replace it with a key in a leaf

Delete the key in the leaf

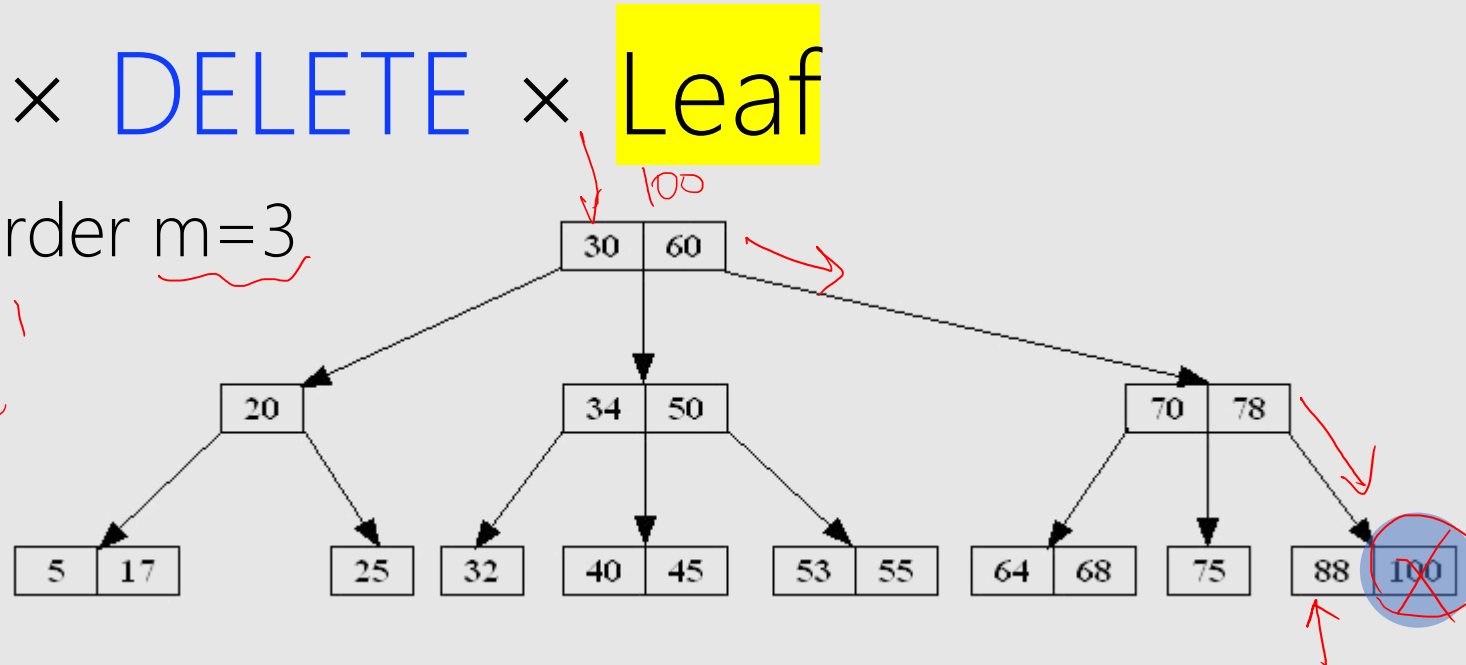
The actual delete always happen at leaf node.

B-tree × DELETE × Leaf

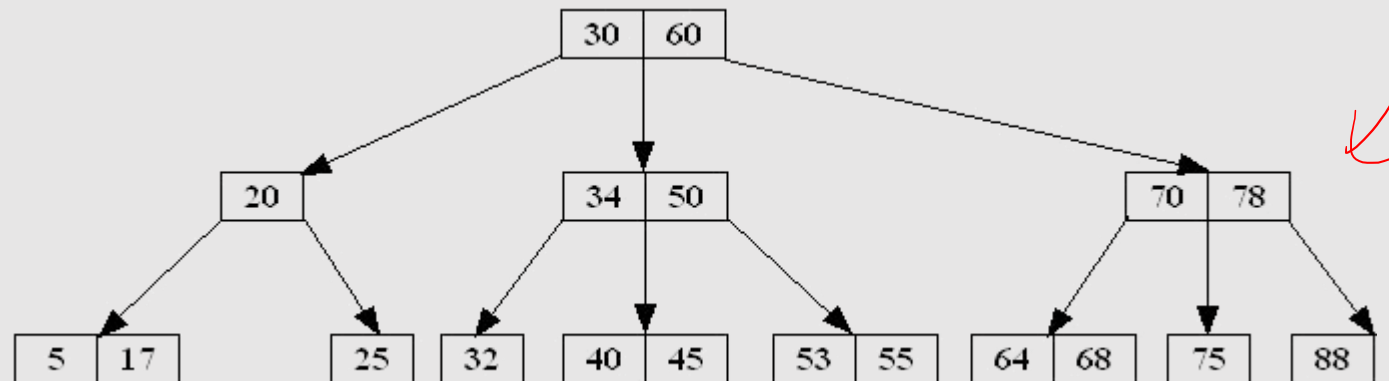
33

B-tree of order $m=3$

$3/2 = 1 \leq 3-1$
 $2 \leq 3-1$



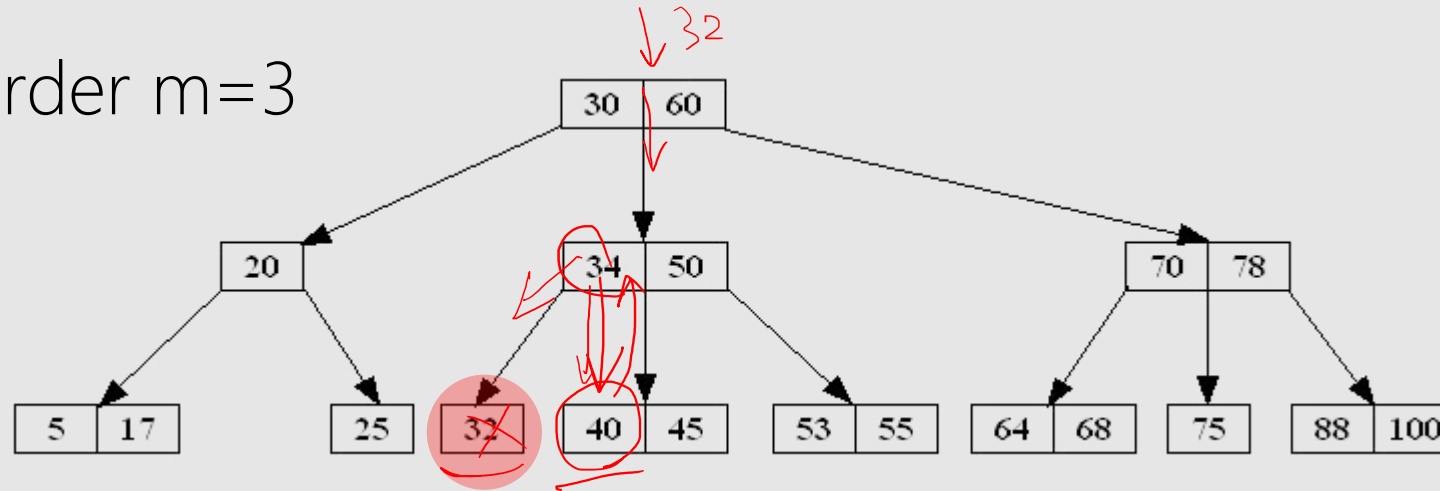
Simple delete key in leaf node. There is no underflow.



B-tree × DELETE × Leaf

34

B-tree of order $m=3$



Delete the key, UNDERFLOW!

Merge with right | left sibling & parent key and sort: [34, 40, 45]

OVERFLOW:

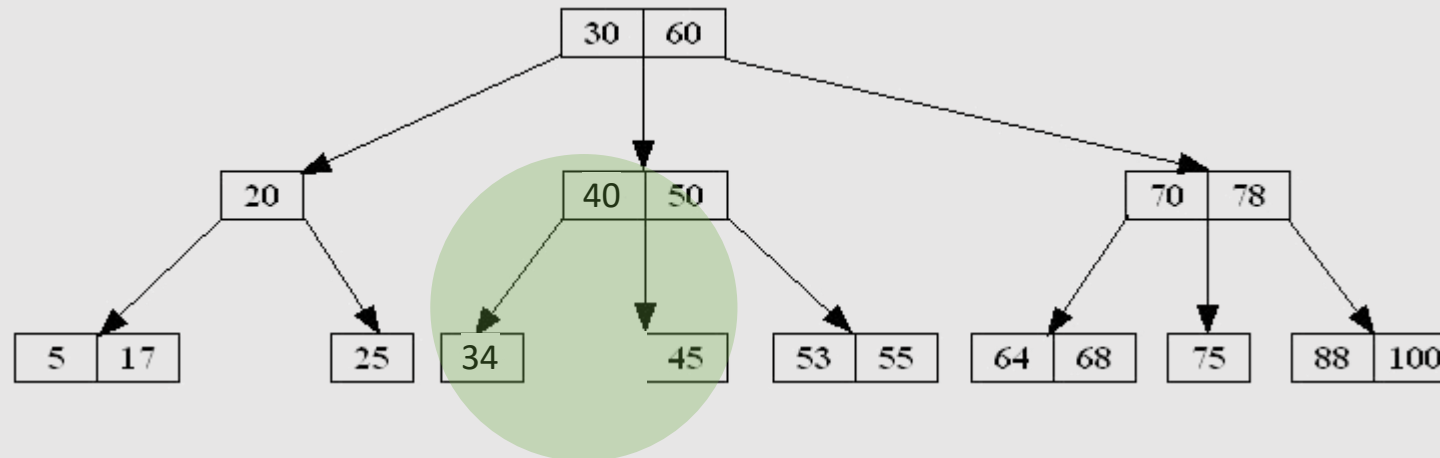
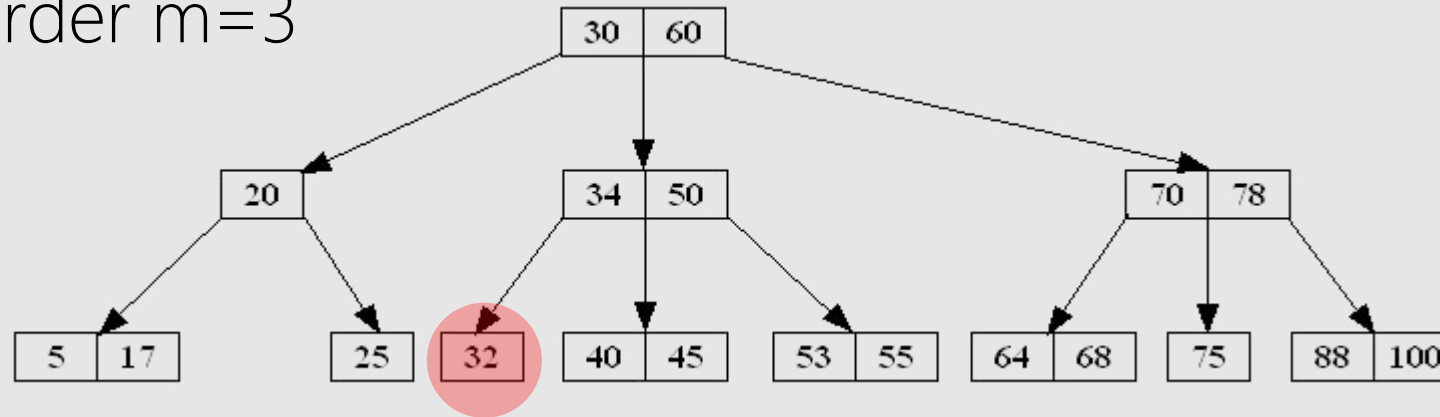
Split by middle key: [34], [45]

Merge middle key to parent: [40]

B-tree × DELETE × Leaf

35

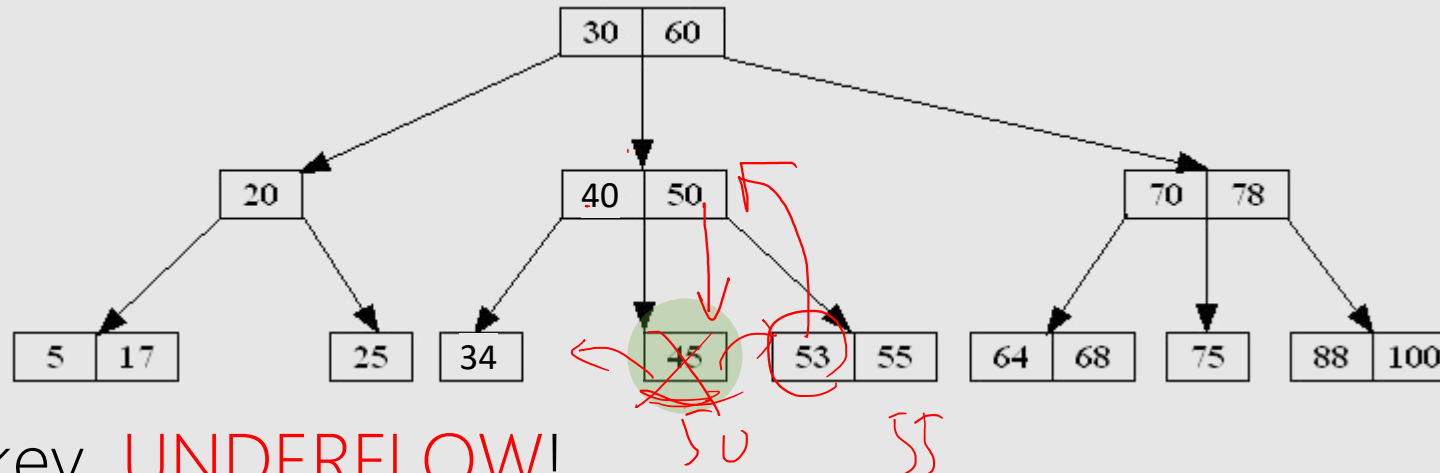
B-tree of order $m=3$



B-tree × DELETE × Leaf

36

B-tree of order $m=3$



Delete the key, **UNDERFLOW!**

Merge with **right** | **left** sibling & parent key and sort: [50, 53, 55]

OVERFLOW:

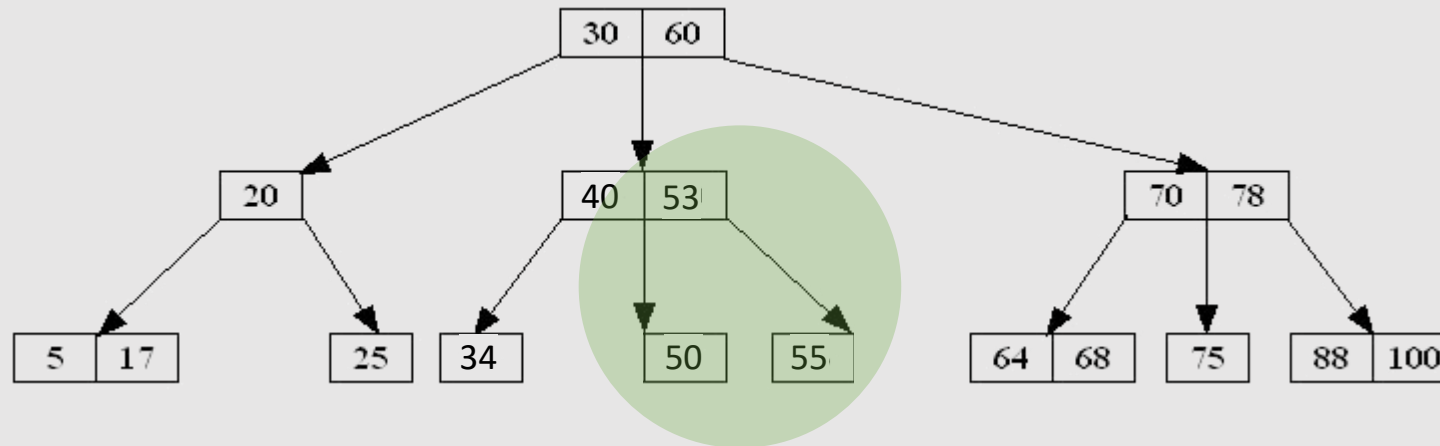
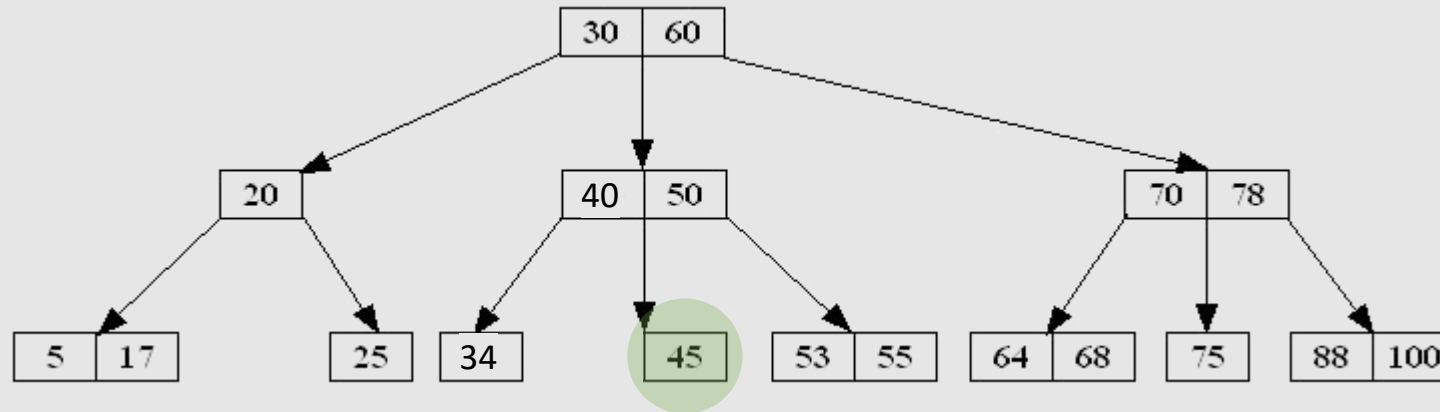
Split by middle key: [50], [55]

Move middle key to parent: [53]

B-tree × DELETE × Leaf

37

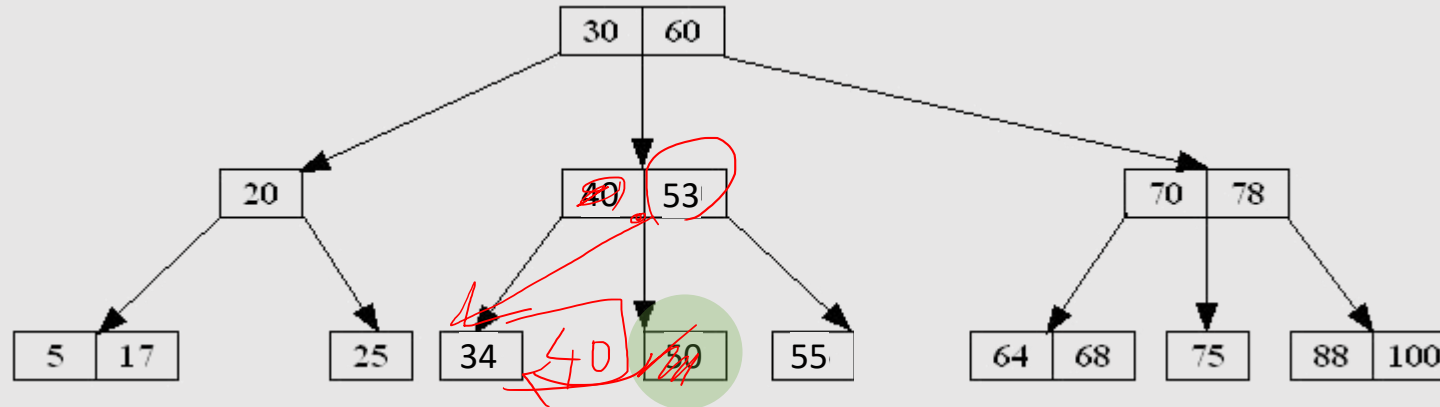
B-tree of order $m=3$



B-tree × DELETE × Leaf

38

B-tree of order $m=3$



Delete the key, **UNDERFLOW!**

Merge with **right** | **left** sibling & parent key and sort: [53, 55]

~~OVERFLOW:~~

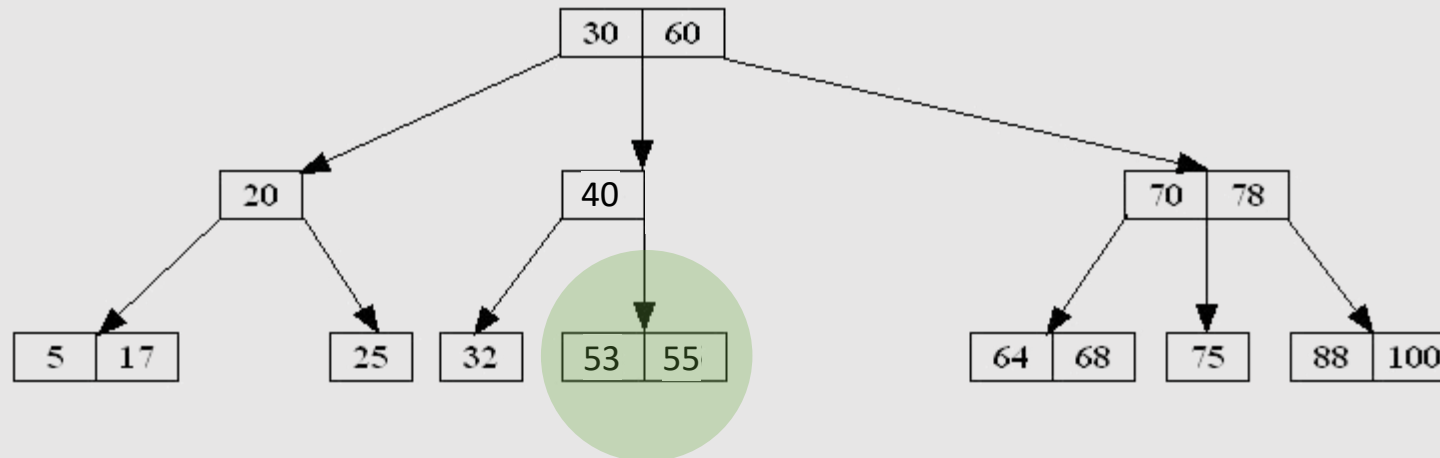
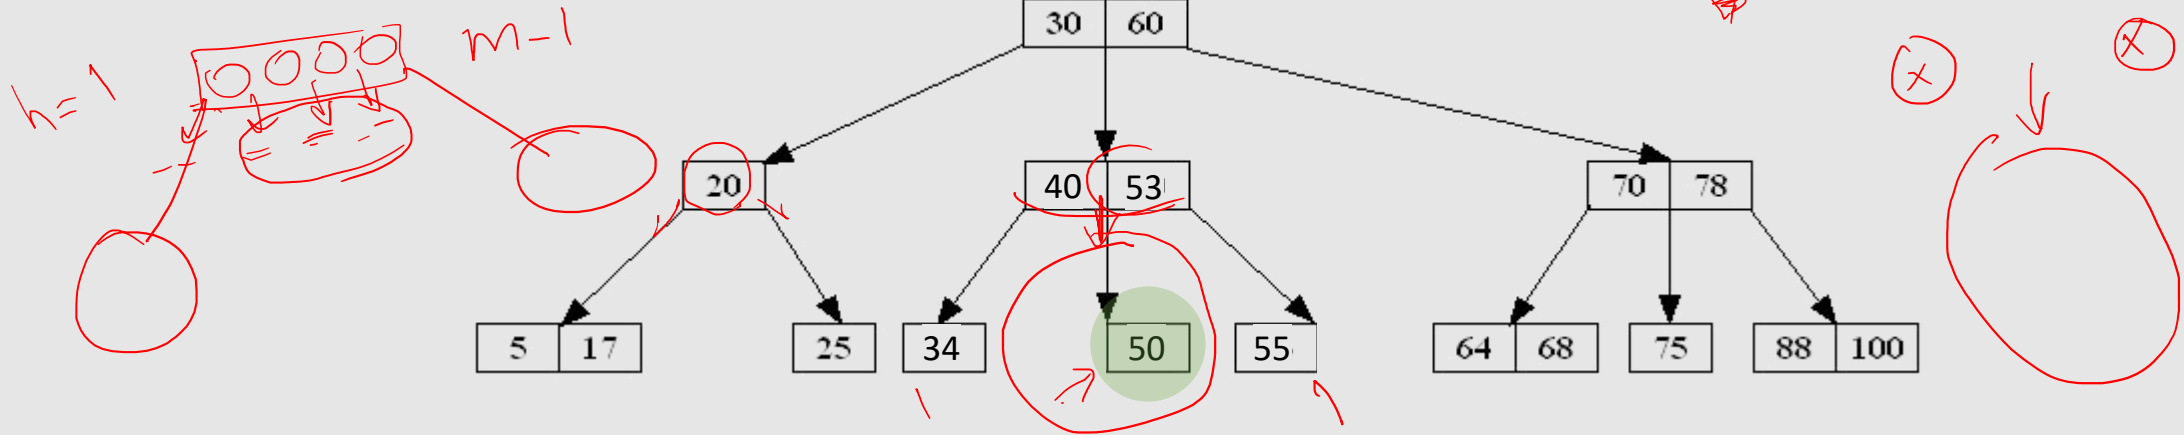
~~Split by middle key~~

~~Merge middle key to parent~~

B-tree × DELETE × Leaf

39

B-tree of order $m=3$



B-tree × DELETE × Leaf

40

DELETE(key):

Find the key

IF the key is in a leaf node

Delete the key

IF (NOT root) & UNDERFLOW

Merge with a adjacent sibling with more keys (right | left) & parent key

IF OVERFLOW

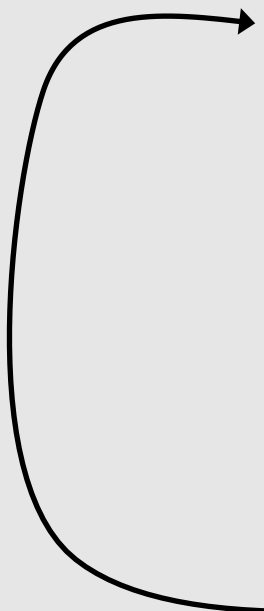
Split based on middle key

Replace middle key with the parent key

ELSE //parent node lost a key and might UNDERFLOW

IF UNDERFLOW in parent

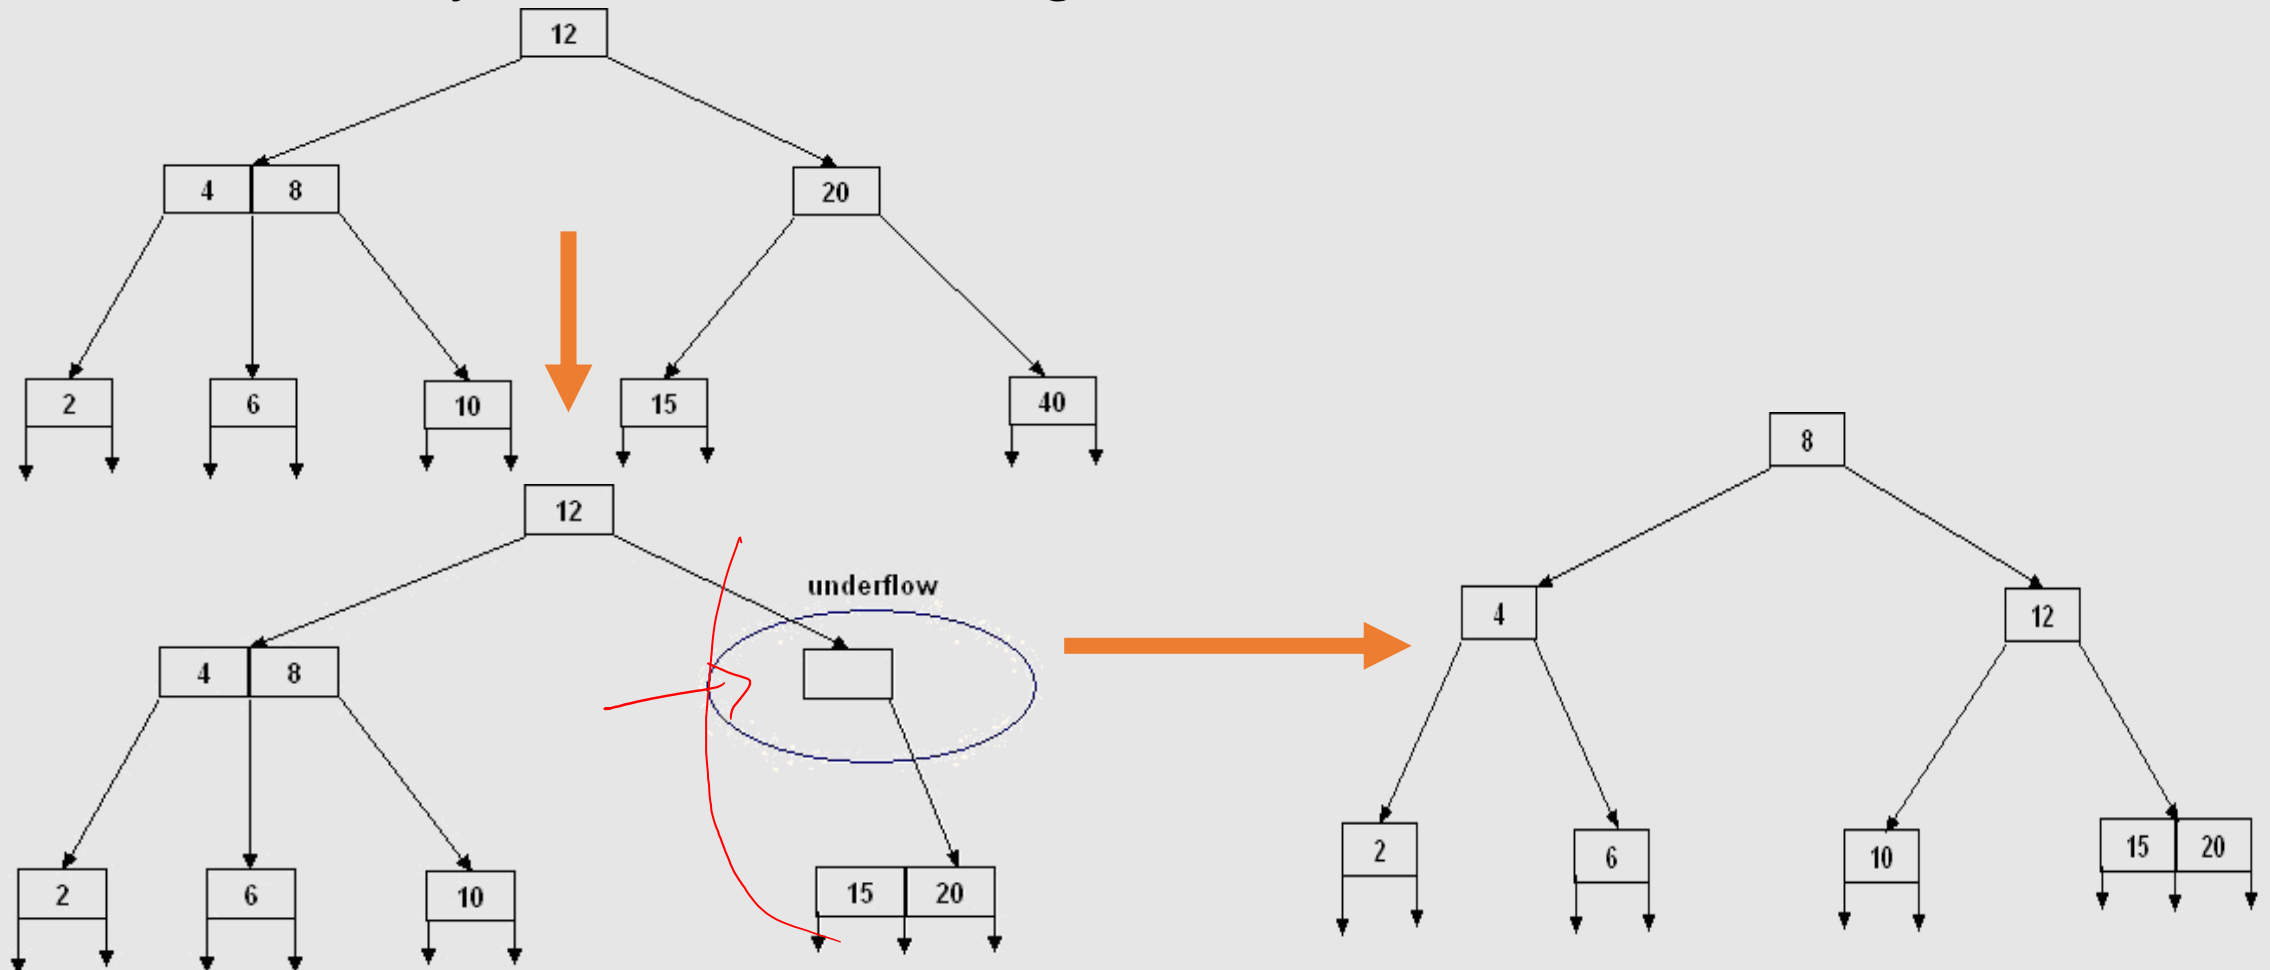
GO TO



B-tree × DELETE × Leaf

40

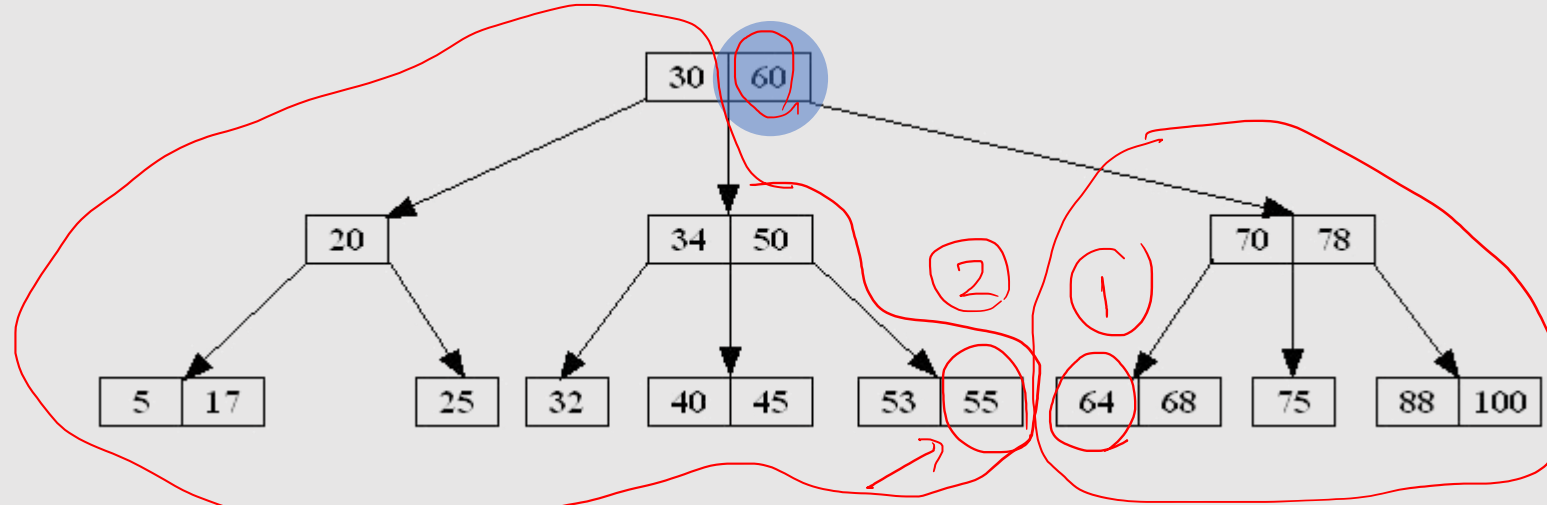
Delete the key 40 in the following B-tree of order 3:



B-tree × DELETE × Non-leaf

41

B-tree of order $m=3$

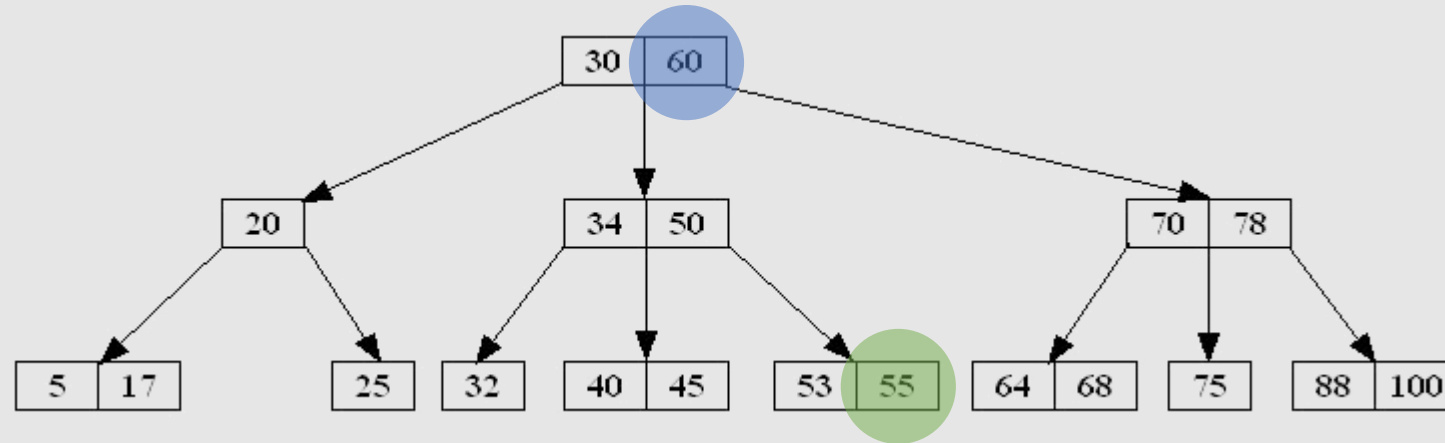


To replace the current key with a key in leaves, what would be the options?

B-tree × DELETE × Non-leaf

42

B-tree of order $m=3$

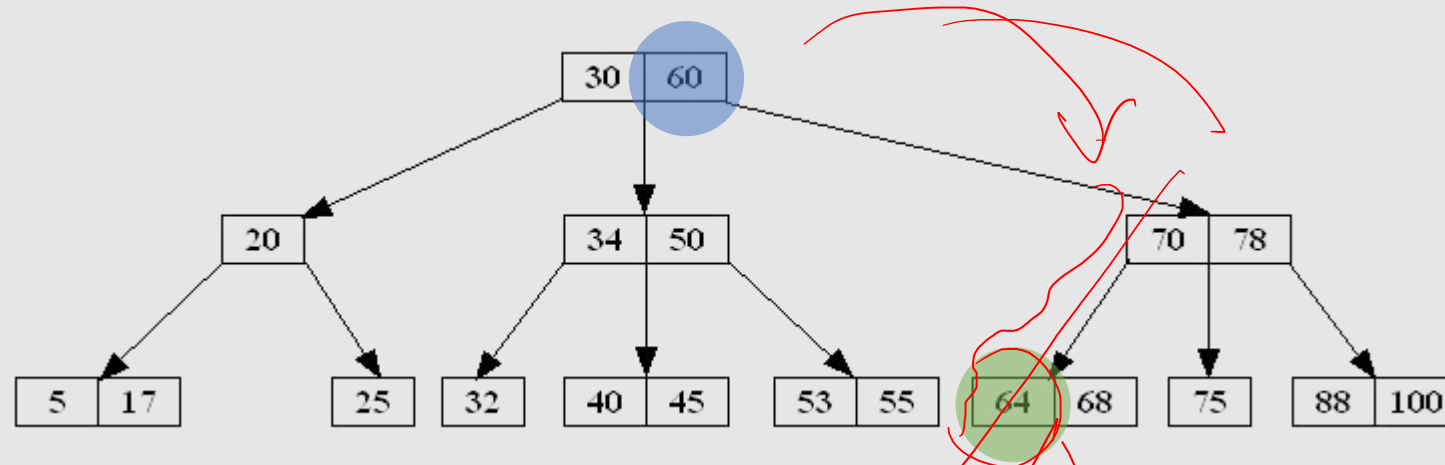


PREDECESSOR(key): Largest key less than the current key
Largest key in left subtree of the current key
Right-most key in left subtree of the current key

B-tree × DELETE × Non-leaf

43

B-tree of order $m=3$

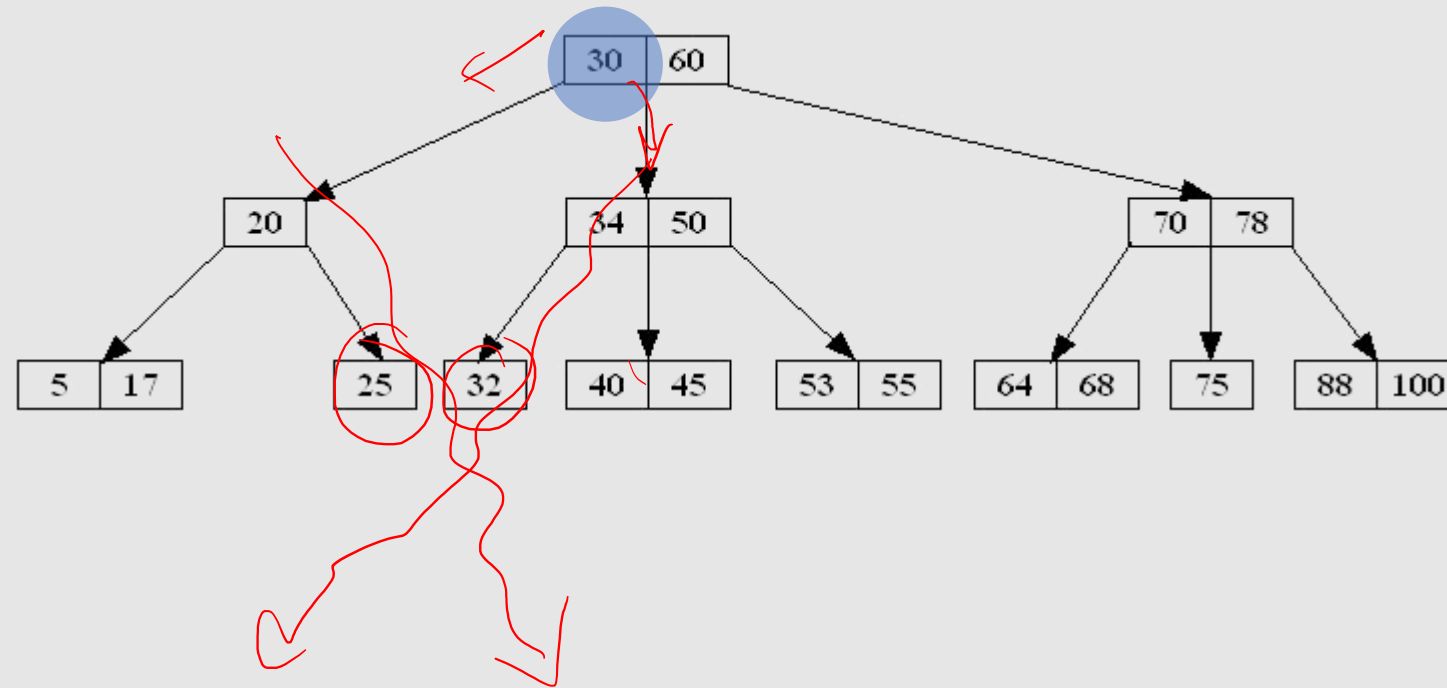


SUCCESSOR(key): Smallest greater than the current key
Smallest in right subtree of the current key
Left-most key in right subtree of the current key

B-tree × DELETE × Non-leaf

44

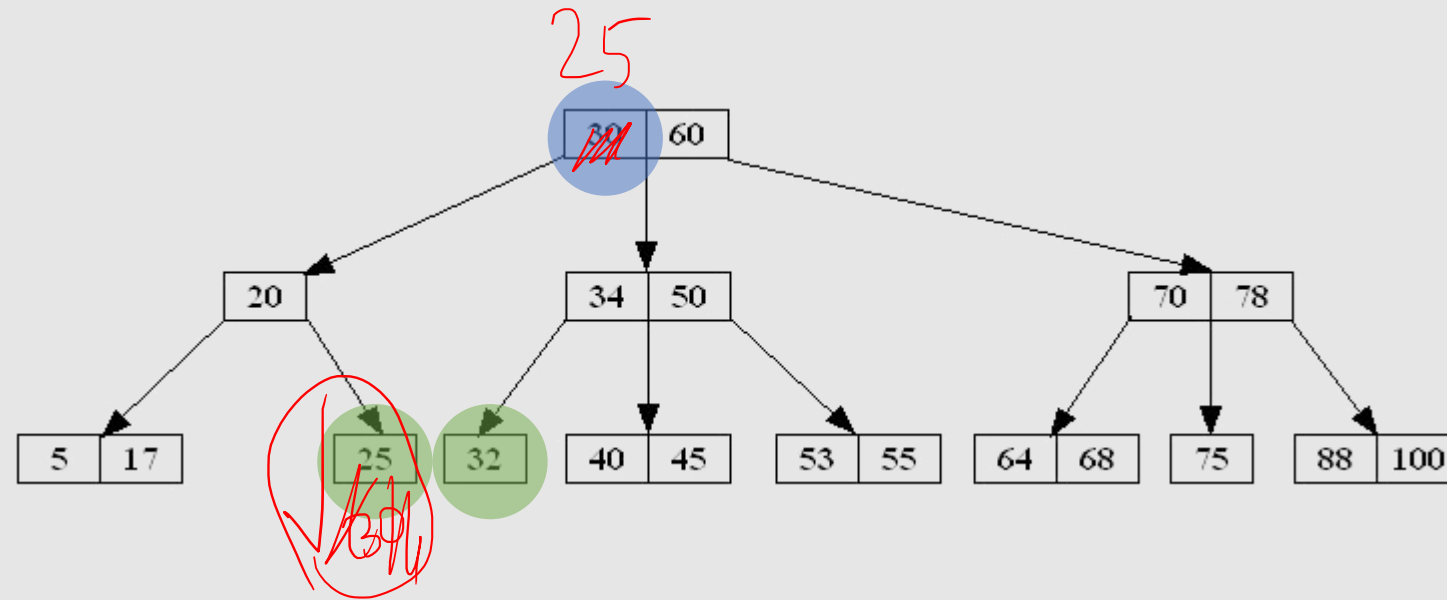
B-tree of order $m=3$



B-tree × DELETE × Non-leaf

45

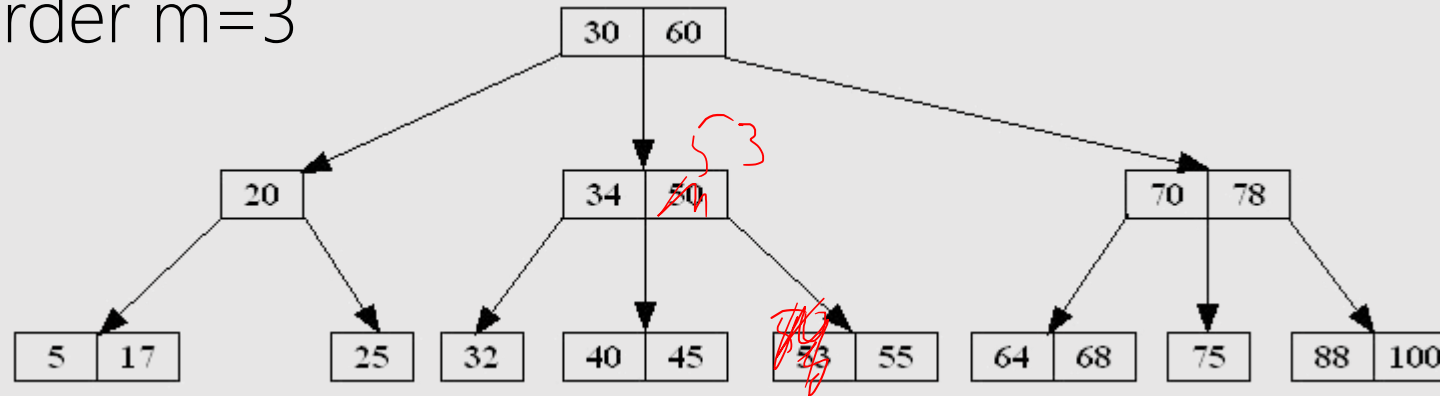
B-tree of order $m=3$



B-tree × DELETE × Non-leaf

46

B-tree of order $m=3$

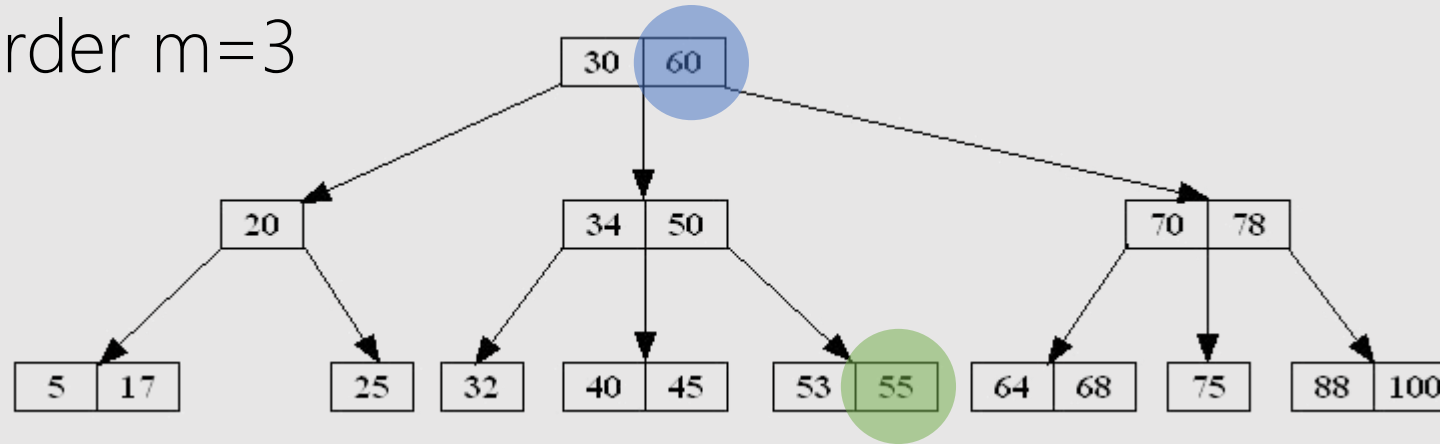


PREDECESSOR	Key	SUCCESSOR
17	20	25
25	30	32
32	34	40
45	50	53
55	60	64
68	70	75
75	78	88

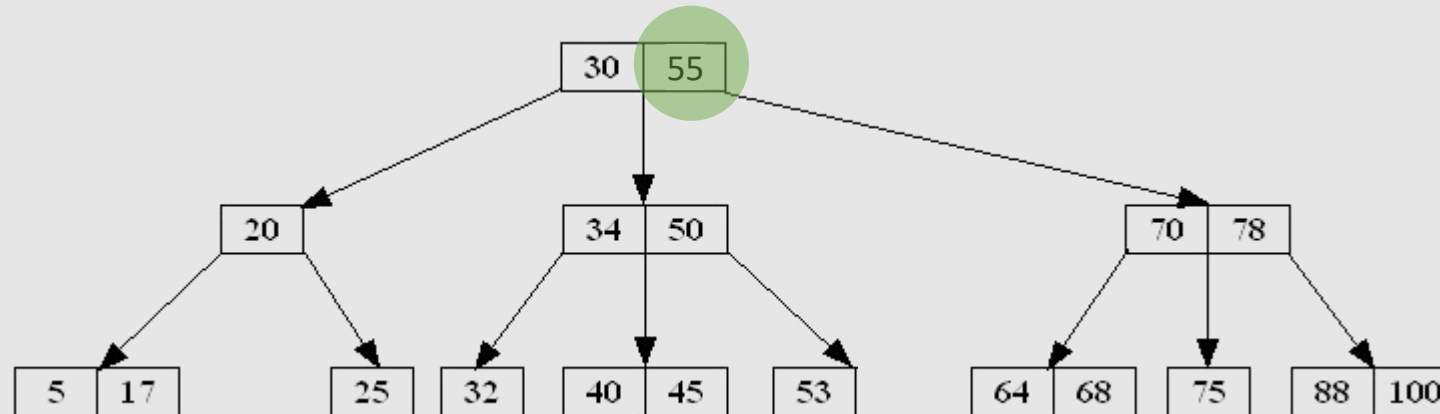
B-tree × DELETE × Non-leaf

47

B-tree of order $m=3$



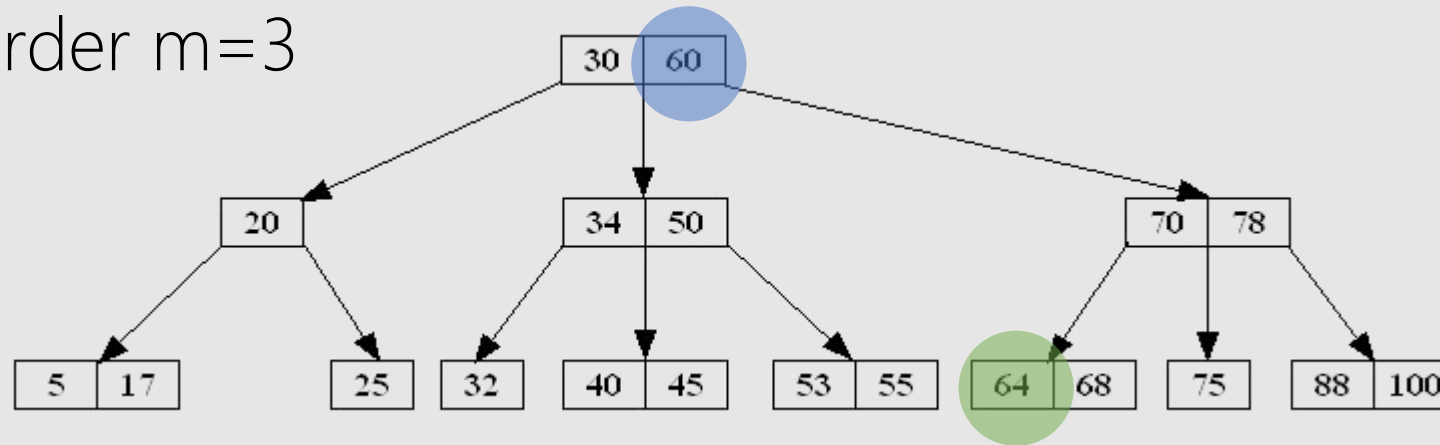
Replace key with PREDECESSOR in non-leaf node



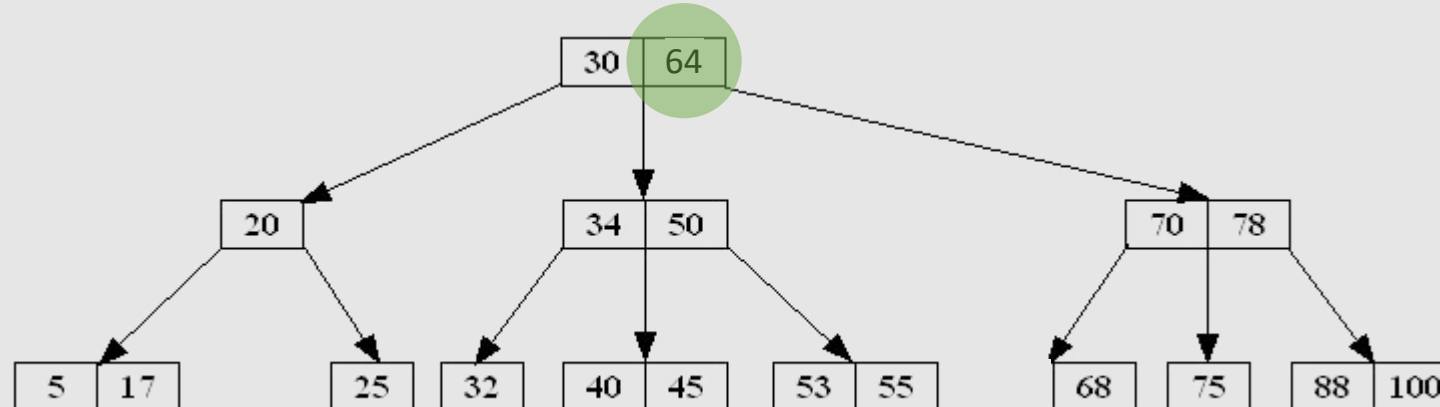
B-tree × DELETE × Non-leaf

48

B-tree of order $m=3$



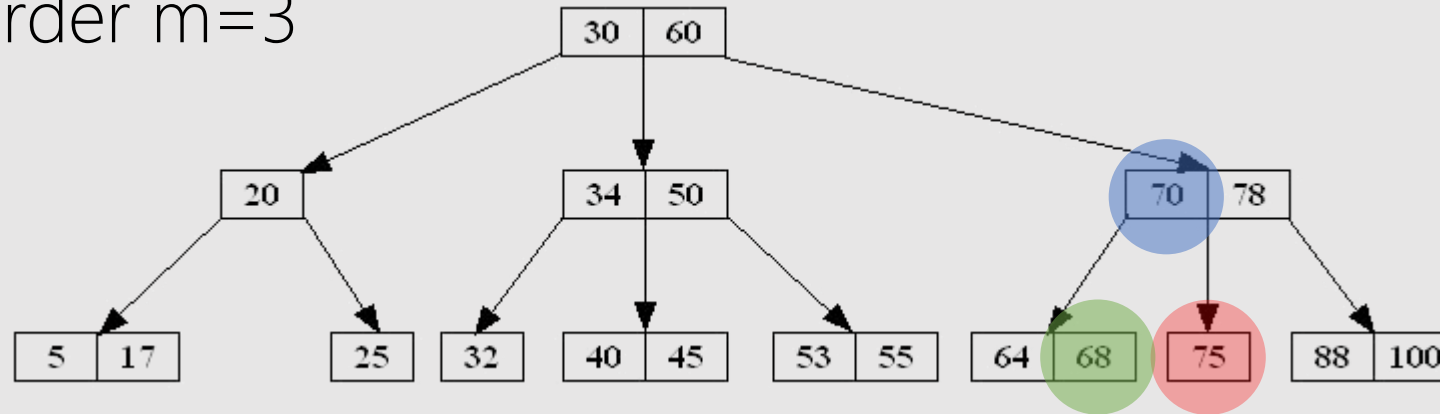
Or replace key with **SUCCESSOR** in non-leaf node



B-tree × DELETE × Non-leaf

49

B-tree of order $m=3$

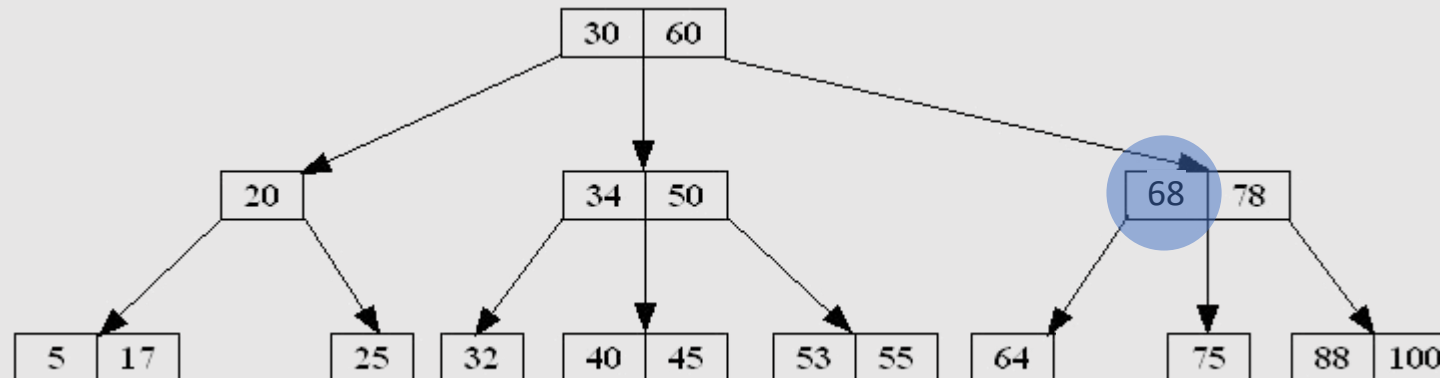
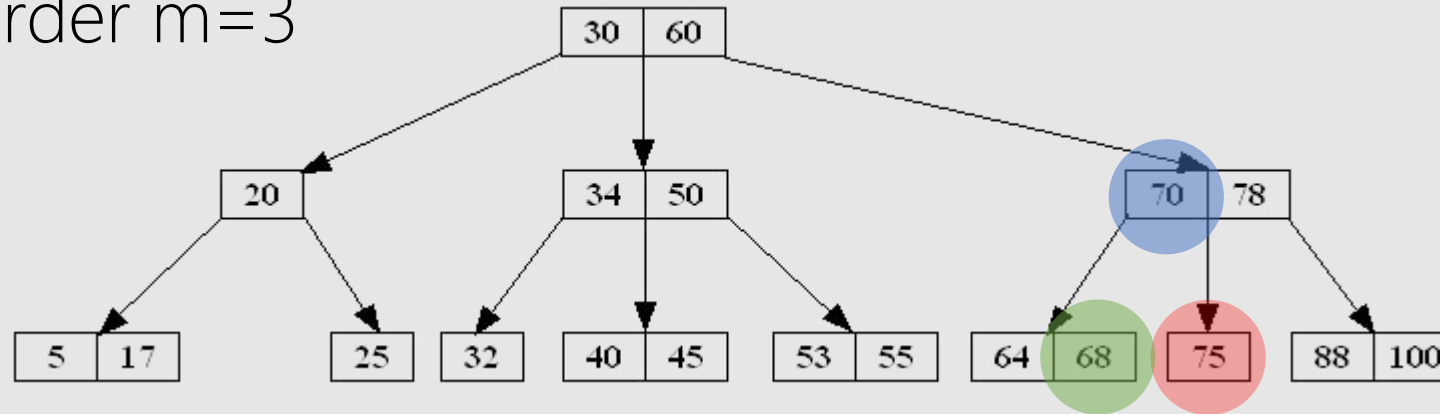


To choose between **SUCCESSOR** | **PREDESSESOR**:
The one whose node is bigger.

B-tree × DELETE × Non-leaf

50

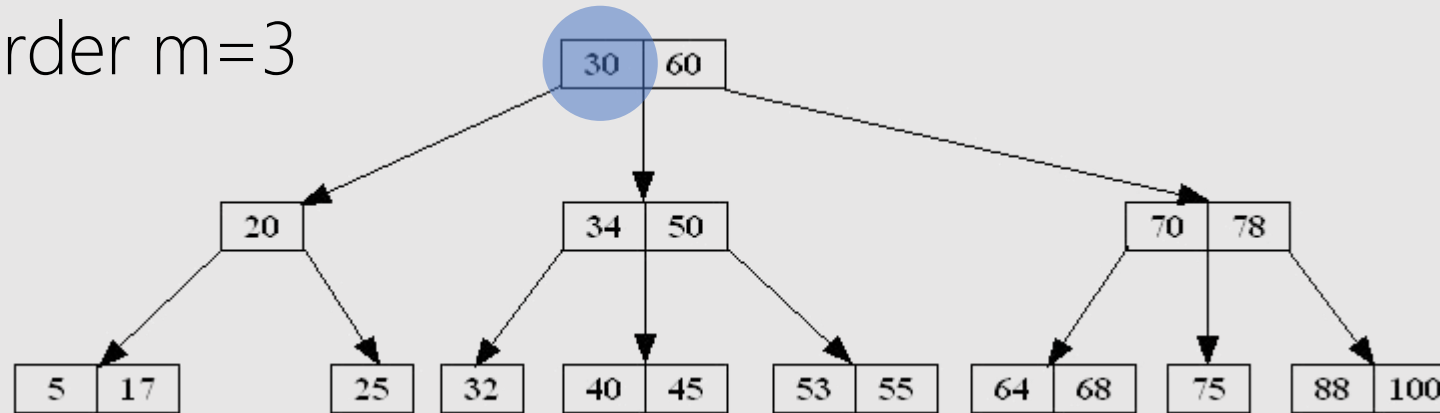
B-tree of order $m=3$



B-tree × DELETE × Non-leaf

51

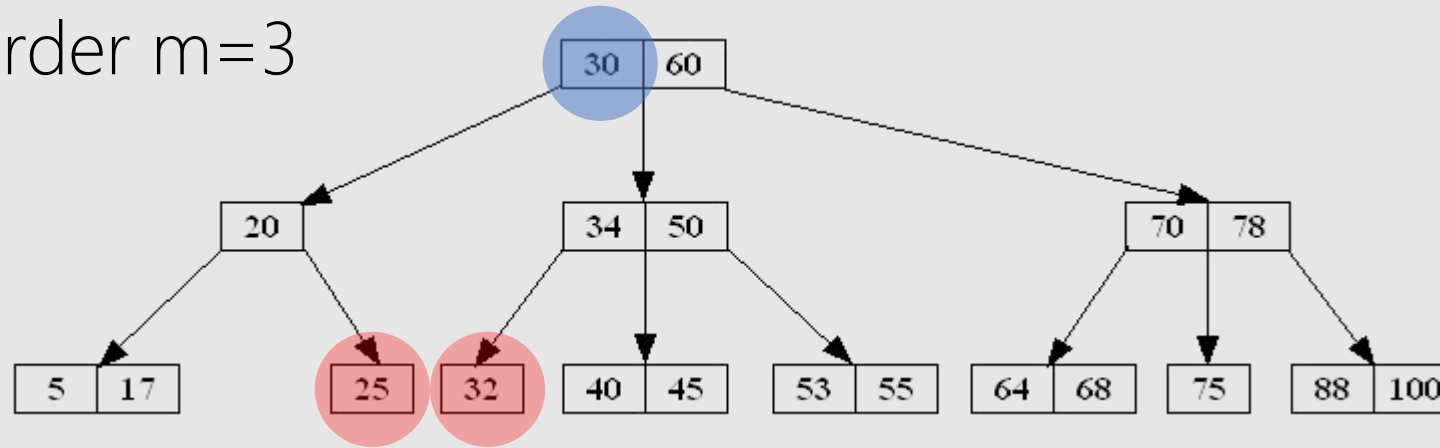
B-tree of order $m=3$



B-tree × DELETE × Non-leaf

52

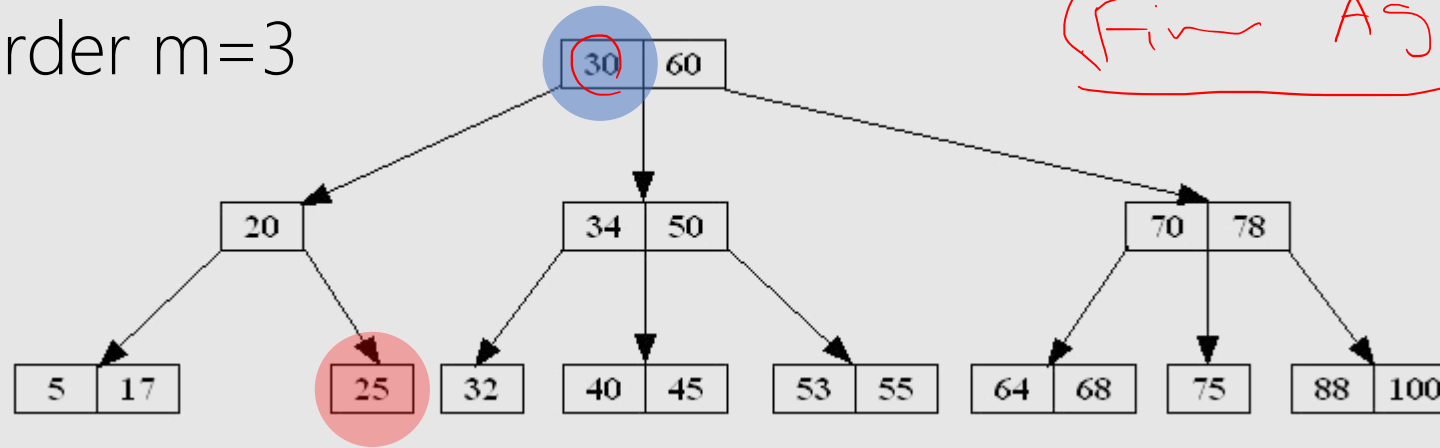
B-tree of order $m=3$



B-tree × DELETE × Non-leaf

53

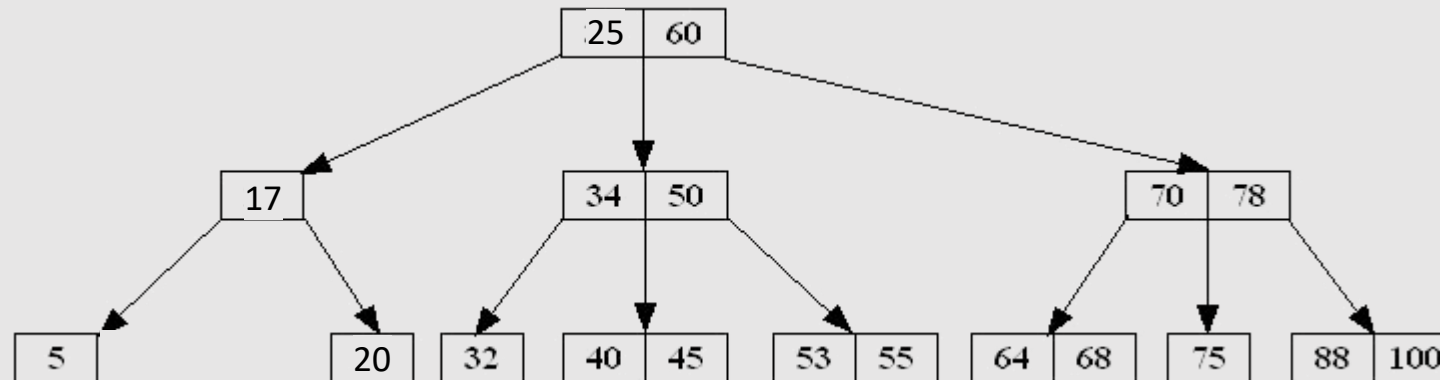
B-tree of order $m=3$



Fin Age

Ali 30
Hoss 40

Alice 30



DBMS × INDEX × B-tree × DELETE

54

DELETE(key):

Find the key

IF the key is in a non-leaf node

key' = PREDECESSOR(key)

key'' = SUCCESSOR(key)

IF |NODE(key')| >= |NODE(key'')|

DELETE(key')

ELSE

DELETE(key'')

The actual delete always happen at leaf node.

DBMS × SQL × INDEX

55

SELECT * FROM Director WHERE LastName = 'Kubrick' AND FirstName = 'Stanley'

Which one?

A) CREATE INDEX IX_LastName_FirstName ON Director(LastName, FirstName)

B) CREATE INDEX IX_FirstName_LastName ON Director(FirstName, LastName)

C) CREATE INDEX IX_FirstName ON Director(FirstName)

D) CREATE INDEX IX_LastName ON Director(LastName)

E) All

F) A & B are the same

50%

50%

DBMS × SQL × INDEX

55

SELECT * FROM Director WHERE LastName = 'Kubrick' AND FirstName = 'Stanley'

Which one?

A) CREATE INDEX IX_LastName_FirstName ON Director(LastName, FirstName)

100% helps

A) CREATE INDEX IX_FirstName_LastName ON Director(FirstName, LastName)

100% helps

A) CREATE INDEX IX_FirstName ON Director(FirstName)

50% helps. Finding FirstName is fast, but after that sequential search to find LastName

A) CREATE INDEX IX_LastName ON Director(LastName)

50% helps. Finding LastName is fast, but after that sequential search to find FirstName

A) All

Having 4 indexes on a table is too much overhead

A) A & B are the same

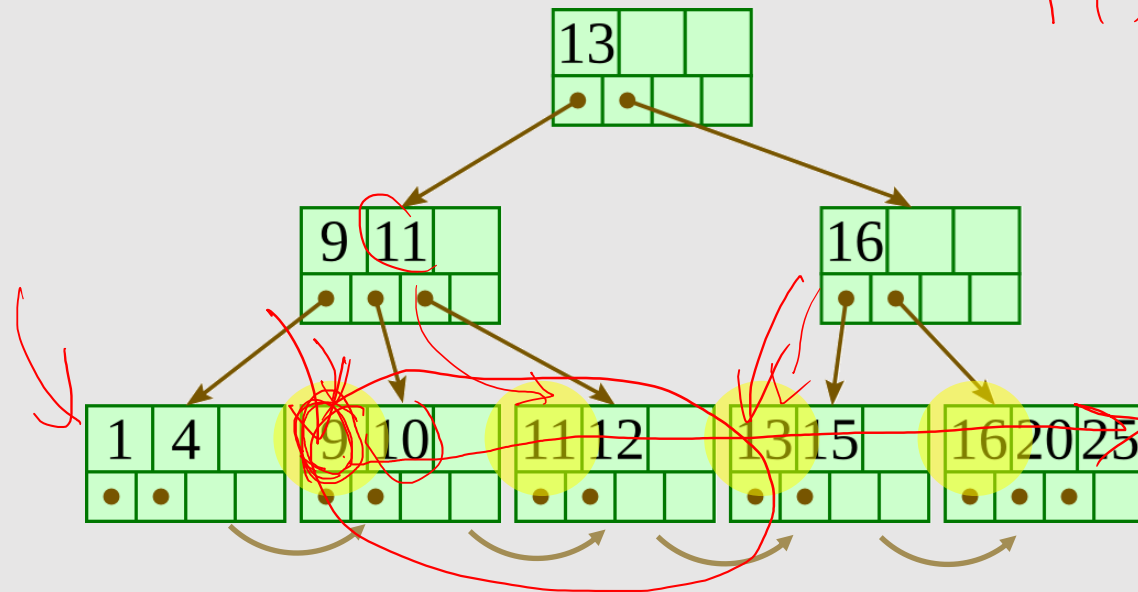
The sorting in B-tree is different, but the effect (help) is the same.

DBMS × INDEX × B⁺-Tree

56

B+-tree is an extension to B-tree in which

- Leaf nodes has a copy of all parents key.
- Leaf nodes are linked together in order to support range queries



DBMS × INDEX × B⁺-Tree

57

B⁺-tree is an extension to B-tree in which

- Leaf nodes has a copy of all parents key.
- Leaf nodes are linked together in order to support range queries

SELECT * FROM Movie

WHERE ReleaseDate BETWEEN 1960 AND 2000

B⁺-tree is the *de facto* standard for indexing in DBMS

DBMS × INDEX × B⁺-Tree

58

CREATE INDEX IX_DateOfBirth_PlaceOfBirth ON Director(DateOfBirth, PlaceOfBirth)

Helps with which one?

- A) SELECT * FROM Director WHERE DateOfBirth = 1955 AND PlaceOfBirth = 'USA'
- B) SELECT * FROM Director WHERE DateOfBirth = 1955
- C) SELECT * FROM Director WHERE PlaceOfBirth = 'USA'
- D) SELECT * FROM Director WHERE DateOfBirth BETWEEN 1955 AND 1980
- E) ALL
- F) None

DBMS × INDEX × Best Practice

59

What columns should be indexed by default on any DB?

Why?

DBMS × INDEX × Best Practice

60

A. Unique Indexes on Primary Keys

`SELECT`s usually have primary keys in `WHERE` clause

Joins are usually based on primary keys equal to foreign keys

B. Unique Indexes on Candidate Keys

To provide uniqueness check

`SELECT`s usually have candidate keys in `WHERE` clause

C. Indexes on Foreign Keys

Joins are usually based on primary keys equal to foreign keys

DBMS × INDEX × Book

61

2nd Ed.

CH08: Indexes in SQL

CH14: Index Structures (B-tree, HashTable, Bitmap, ...)

1st Ed.

CH06: The Database Language SQL (6.6.5 & 6.6.6)

CH13: Index Structures (B-tree, HashTable)

CH14: Multidimensional & Bitmap Indexes